

# PRINCIPLES OF COMPUTER SYSTEMS DESIGN

## CSE130

Winter 2020

### History



### Notices

- **Administration 1 & 2** due 23:59 **Wednesday January 15**
- **Lab 1** due 23:59 **Sunday January 19 CHANGED!**
- **NO LECTURE Friday January 10**

2

### Today's lecture

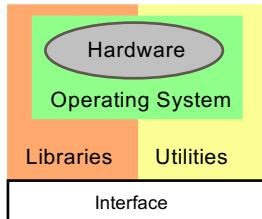
- Brief History of Computer Systems
- Structure of the course
- Lab 1 "Secret Sauce"

### What is a Modern Operating System?

- The GUI is not the OS, it's just the visible part
- We are interested in the basic engineering principles underlying the most used computer systems
- **The OS is the resource manager for a computer**
  - Controls and arbitrates access to hardware:
    - CPU
    - Memory
    - I/O Devices
  - It provides high level abstractions to application programmers for hardware/device control
  - It provides services for multiprogramming (usually)

## Basic Computer System Architecture

- Hardware:
  - CPU
  - Memory
  - I/O Devices
- Operating System:
  - Coordinates use of resources amongst applications and users
- Libraries and Utilities:
  - Provide standard functionality



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

5

## Operating System goals

- **Efficient use of Resources**
  - Maximize utilization of resources
  - Avoid bottlenecks
- **User Convenience and Productivity**
  - User time is expensive
- **Availability and Reliability**
  - Many systems are critical
  - Data is valuable

Efficiency  
Convenience  
Availability  
Reliability

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

6

## Design Tradeoffs

- Cannot simultaneously optimize **efficiency** and **user convenience**
- Choice depends on context...

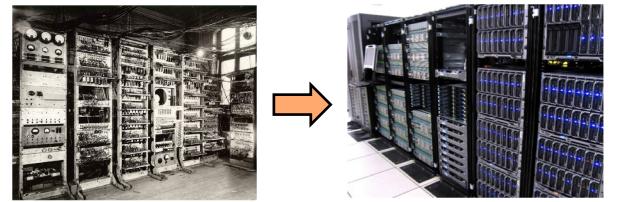


UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

7

## Evolution of Computer System

- 70+ years of development
  - Simple, single-purpose systems to all-singing, all-dancing hand-held devices, via general-purpose super computers



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

8

2

## Batching Systems

- With early machines the operators had to manually set up the runtime context for each program - a labor intensive task
  - To reduce the load on the (human) operator, jobs were sorted into batches with similar runtime requirements
  - Setting up the environment for a new batch required considerable effort and time so the human operators made little devices or wrote little utilities to make life easier for themselves
- Operating Systems started as batch job control systems**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

9

## Batching Systems

- These original Operating Systems were simple control programs that automatically worked through individual jobs in a batch (sequence of jobs)



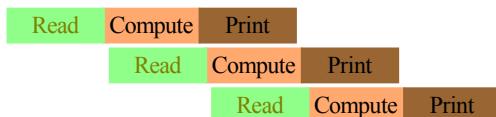
- IO (read/print) slow, CPU fast(er) => under utilization

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

10

## Spooled Batching

- Next increment was spooling batch systems
  - This is the same idea as CPU instruction pipelining



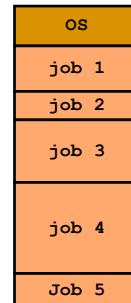
- Requires hardware interrupts ( more on these later )
- IO still slow, CPU still fast(er) => a bit better

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

11

## Multi-Programmed Batching

- Multi-programming allows multiple jobs (programs) to be resident (in memory) at the same time, all sharing a single CPU
- When one program is blocked on IO, we simply pick another and execute it
- As long as at least one job remains ready to execute, the CPU is 'fully' utilized



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

12

## Multi-Programmed Batching

- The start of sophisticated Operating Systems
  - Job scheduling (which job in batch is CPU resident)
  - Memory management (and protection)
  - Requires CPU scheduling (fairness between jobs)
- Still batch though, with human controllers (operators) and no direct user interaction
- Very efficient, high hardware utilization

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

13

## The Stretch

- In November 1956, IBM set out to build a computer 100 x faster than its current 704 model - a lofty goal of 4 million instructions a second (MIPS)
- The 7030 was IBM's first transistorized computer and pioneered many advanced computing concepts:
  - **Interrupts**
  - Memory Interleaving
  - **Memory Protection**
  - Memory Error Detection and Correction
  - **Multiprogramming**
  - Pipelining
  - Immediate operands
  - Instruction pre-fetch
  - Speculative execution, result forwarding
- Only achieved ~30x the 704 (1.2 MIPS)
- Fastest computer in the world for 4 years
- First machine shipped to Los Alamos in 1961



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

14

## Time Sharing Systems

- Natural evolution from multi-programmed systems
- Add user interaction
  - Switch between user programs frequently, so users can interact with each program
  - $N$  people, so each person effectively has  $1/N$  of the computer's time
    - **Or do they?**
  - Many users sharing same machine - IO blocks on user response and program instances are switched by the OS
- Resulted in a significant change in OS complexity
  - There were many, many failures in the development of time sharing systems  $\text{\_}\text{\_}\text{\_}\text{\_}\text{\_}$
- **Time Sharing  $\Rightarrow$  Multiprogramming**
- **Multiprogramming  $\not\Rightarrow$  Time Sharing**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

15

## The CDC 6600

- Considered the first supercomputer, delivered 1964
- Fastest computer in the world until 1969
  - 10 MIPS @ 10 MHz
- 10 superscalar functional units in the CPU
- Introduced a new OS
  - Simultaneous Processing Operating System
- SPOS was a failure  $\ominus$ 
  - Reverted to a simple Job Control



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

16

## Cray Supercomputers

- CDC 7600 Replaced the CDC 6600 in 1969
- Fastest Super Computer until 1975
  - 10 to 36 MIPS @ 36 MHz
- First functional timesharing OS
  - Livermore Time Sharing System (LTSS)



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.



- Seymour Cray left CDC in 1972 and formed Cray Research
- Cray-1 launched in 1976
  - 160 MIPS @ 80 MHz
- Cray Time Sharing System (CTSS) derived from LTSS on the CDC 7600

17

## Personal Computers



Apple 1 (1976)

Commodore Pet (1977)

IBM 5150 (1981)

18



oldcomputr.com

Sharp MZ80K (1978)

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

19

## Personal Computers

- Paradigm shift
  - The user is now most important
  - Hardware is low cost
- Goal is convenience and responsiveness, not hardware utilization
- Did not initially perform multitasking or handle multiple simultaneously connected users
- As power increased, these (and many other) features migrated into the mainstream
- Modern PCs now provide multiple processing units in a shared memory architecture

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

20

## Cluster Computing

- Many cheap commodity computers
- Clustered: Linked with a high speed LAN
- MIMD Parallelism
  - Multi Instruction, Multi Data
- Beowulf cluster – Linux
- Can be for availability, performance, or both (load balanced)



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

21

## Distributed Computer Systems

- A loosely coupled collection of processors:
  - Do not share clock or memory
  - Do share a communications network
  - Heterogeneous
  - Permit sharing of resources (CPU, files, ...)
  - Issues include:
    - load balancing
    - process migration
    - distributed transactions
    - security
    - communications
    - reliability



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

22

## Structure of the course

- Five areas:
  - Processes, Threads, and Concurrency
  - CPU Scheduling
  - Memory Management
  - File Systems and Mass Storage
  - Miscellaneous
    - Virtualisation, Security, History, etc.
- Final examination has five questions, one for each area

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

23

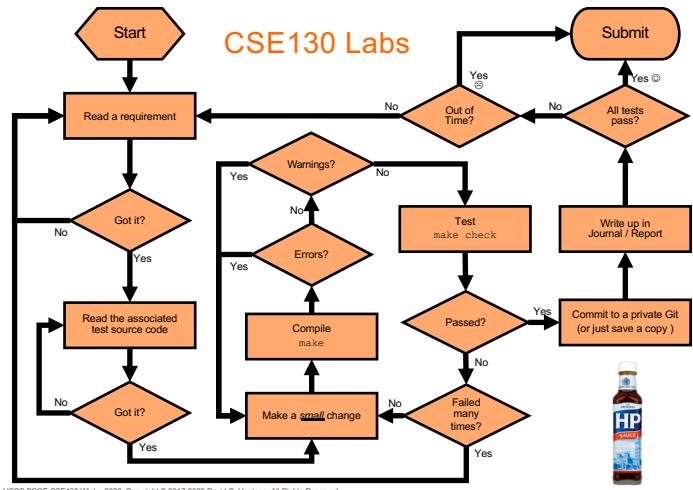
## What is Secret Sauce?

- General advice
- Extra detail on labs
- Video explanations
- Step-by-step guides
- Pseudo code
- After-the-fact feedback



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

24



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

25

## What do you need to do for Lab 1?

- In the `timer_sleep()` function:
  - Instead of busy-waiting, put the current thread to sleep and immediately return
- At regular points in the future:
  - Wake up sleeping threads at or past their wakeup time

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

26

## The `timer_wait()` function

```
/*
 * Sleeps for approximately TICKS timer ticks.
 * Interrupts must be turned on.
 */
void
timer_sleep(int64_t ticks)
{
    int64_t start = timer_ticks();

    ASSERT(intr_get_level() == INTR_ON);

    while (timer_elapsed(start) < ticks)
        thread_yield();
}
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

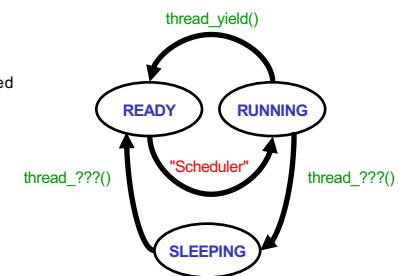
28

## Pintos Thread States & Transitions

**READY**  
Waiting to run on the CPU

**RUNNING**  
Binary instructions are being executed

**SLEEPING**  
Waiting until it's time to wakeup



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

29

## Questions to ask yourself (not exhaustive)

- About Pintos

- Where are timer functions defined / implemented?
- What is the C type of a thread?
- How do I get a reference to the current thread?
- How do I put a thread to sleep?
- How do I wake a thread up again?
- When do I check if anything needs waking up?
  - i.e. what function gets called when the timer interrupt fires? (massive hint)
- How do I turn interrupts on and off?

- About your solution

- How might I store references to threads I want to wake?
- How do I store the wakeup times?
- Would my solution handle 1,000's of threads?

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

30

## Next Lecture

- **NO LECTURE Friday January 10**

- Monday: Processes & Threads I

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

38