

# PRINCIPLES OF COMPUTER SYSTEMS DESIGN

## CSE130

Winter 2020

Concurrency I



### Notices

- **Lab 1** due 23:59 **Sunday January 19**
- **Assignment 1** due 23:59 **Sunday January 26**
- No Class Monday - Martin Luther King Jr. Day
- Change of Lecture Location

### Baskin Engineering 101

From Wednesday January 22

2

### Today's Lecture

- Processes & Threads Re-Cap
- Thread Models & Pools
- Shared Memory
- Introduction to Concurrency
- Introduction to Assignment 1

### Process States & Transitions

#### NEW:

The process (P) is being created

#### READY:

P is waiting to run on the CPU

#### RUNNING:

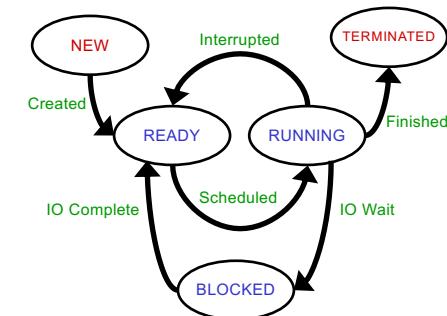
P's instructions are being executed

#### BLOCKED:

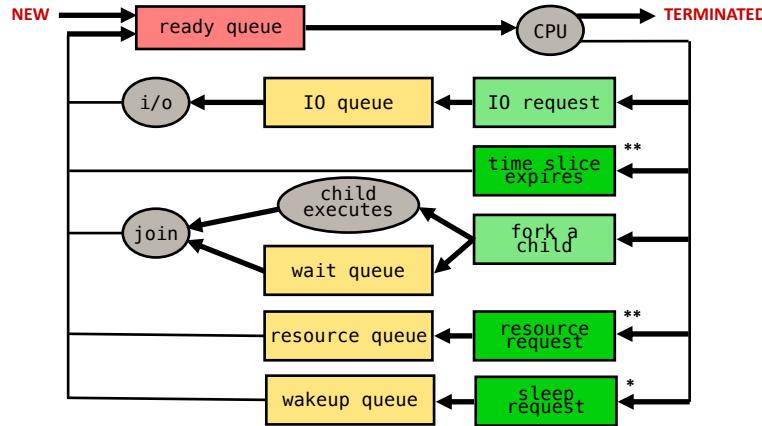
P is waiting on IO to complete

#### TERMINATED:

P's execution is complete



## Process Execution



UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

\* Lab 1   \*\* Lab 2

5

## Recap: Processes & Threads

- **Processes:** “Heavyweight”
  - Expensive to create/switch/destroy
  - Owns resources (files, network sockets, etc.)
- **Kernel-level Threads:** “Lightweight”
  - At least one per process
  - If multiple, they share memory and resources
  - Has own stack, registers, program counter, id
  - Cheaper to create/switch/terminate than process
  - OS can schedule one to each logical core and can swap those that block
    - Super awesome feature, a justification for kernel-level threads all on its own ☺
- **User-level Threads:** “Inexpensive”
  - Implemented by user space libraries
  - Kernel is unaware of them
  - Very fast to create and manage
  - Blocked if the kernel thread (or process) running them is blocked
  - Modern libraries base user-level threads on kernel-level threads to benefit from multiprocessor machines

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

6

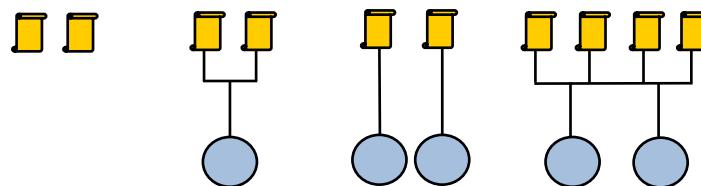
## The Totally Tricky Thread Taxonomy

- **Kernel-level Thread**
  - Thread class managed by the operating system
  - Creation requires a system call unless created by the kernel itself
- **User-level Thread**
  - Thread class managed by a library external to operating system
  - Creation does not necessarily require a system call
  - Kernel is unaware of their existence
- **Kernel thread**
  - Thread instance created by the operating system kernel
  - Implicitly a Kernel-level Thread
- **User thread**
  - Thread instance created by a user process (code written / run by a user)
  - Can be a Kernel-level Thread or a User-level Thread

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

7

## Thread Models (i.e. Mapping User-level Threads to Kernel-level Threads)



**Many-to-None**  
Single process “Stand Alone” user-level threads - kernel is unaware of their existence

**Many-to-One**  
Equivalent to single process user-level threads, and has the same advantages and shortcomings

**One-to-One**  
Expensive, as one kernel-level thread is needed for each user-level thread

**Many-to-Many**  
Best model - neither too slow, nor blocking, but difficult to program

User-level thread Kernel-level Thread

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

8

## Thread Pools

- If a (web or similar) server creates a new thread to serve each client:
  - Costs time to set up the new thread
  - Thread is thrown away afterwards
  - Wasteful of resources
  - Difficult to manage due to resource limits
- Instead, such a server can create a **pool of threads** to serve requests - threads are recycled back to the pool on completion
- This is easy to manage and implement, and only consumes a small amount of additional runtime resource
- Care must be taken to expunge state when the threads are recycled
  - **Why?**
    - Security

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

9

## Inter Process Communication (IPC)

- Consider two processes on a Single CPU system
- How do they communicate with each other?
  - i.e. how do they share data?

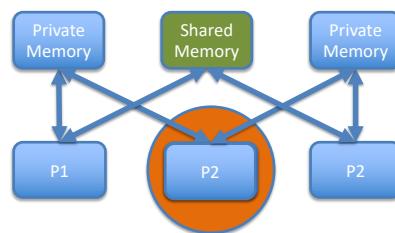


UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

10

## Shared Memory

- Consider two processes on a Single CPU system



UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

11

## Unix / Linux Shared Memory System Calls

- `ftok` Generate a unique key
- `shmget` Get an identifier for a shared memory segment
- `shmat` Attach to the shared memory segment
- `shmdt` Detach from the shared memory segment
- `shmctl` Control (e.g. delete) the shared memory segment

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

<http://man7.org/linux/man-pages/man2/shmget.2.html>

12

```

#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdio.h>
#include <string.h>

int main()
{
    // ftok to generate unique key key_t
    key = ftok("shmseg",65);

    // shmget returns an identifier in shmid
    int shmid = shmget(key,1024,0666|IPC_CREAT);

    // shmat to attach to shared memory
    char *str = (char*) shmat(shmid,(void*)0,0);

    // copy some data into shared memory
    strcat(str, "Hello CSE130!");
    printf("Written to shared memory\n");

    // detach from shared memory
    shmdt(str);

    return 0;
}

```

```

$ ./producer
Written to shared memory

```

```

#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdio.h>

int main()
{
    // ftok to generate unique key
    key_t key = ftok("shmseg",65);

    // shmget returns an identifier in shmid
    int shmid = shmget(key,1024,0666|IPC_CREAT);

    // shmat to attach to shared memory
    char *str = (char*) shmat(shmid,(void*)0,0)
    printf("Data read: %s\n",str);

    // detach from shared memory
    shmdt(str);

    // destroy the shared memory
    shmctl(shmid,IPC_RMID,NULL);

    return 0;
}

```

```

$ ./consumer
Data read: Hello CSE130!

```

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

13

```

#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdio.h>
#include <string.h>

int main()
{
    // ftok to generate unique key key_t
    key = ftok("shmseg",65);

    // shmget returns an identifier in shmid
    int shmid = shmget(key,1024,0666|IPC_CREAT);

    // shmat to attach to shared memory
    char *str = (char*) shmat(shmid,(void*)0,0);

    // copy some data into shared memory
    strcat(str, "Hello CSE130!");
    printf("Written to shared memory\n");

    // detach from shared memory
    shmdt(str);

    return 0;
}

```

```

$ ./producer
Written to shared memory
$ ./producer
Written to shared memory
$ ./consumer
Data read: Hello CSE130!Hello CSE130!

// shmat to attach to shared memory
char *str = (char*) shmat(shmid,(void*)0,0)
printf("Data read: %s\n",str);

// detach from shared memory
shmdt(str);

// destroy the shared memory
shmctl(shmid,IPC_RMID,NULL);

return 0;
}

```

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

14

## POSIX Shared memory System Calls

- `shm_open` Create and open a new shared memory object or open an existing object
- `ftruncate` Set the size of the shared memory object (a newly created shared memory object has a length of zero)
- `mmap` Map the shared memory object into the virtual address space of the calling process
- `munmap` Unmap the shared memory object from the virtual address space of the calling process
- `shm_unlink` Remove a shared memory object by name
- `close` Close the file descriptor allocated by `shm_open` when it is no longer needed

## Concurrency

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

[http://man7.org/linux/man-pages/man7/shm\\_overview.7.html](http://man7.org/linux/man-pages/man7/shm_overview.7.html)

15

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

16

## Concurrency in Computer Systems

- ***The actual or apparent simultaneous execution of threads***
  - Actual = multicore and/or multiprocessors
  - Apparent = time slicing
  - Today's Reality = both
- Threads running concurrently in the same process typically want to access shared data
- Unfortunately, sharing data in an uncontrolled fashion allows that shared data to become:
  - Inconsistent
  - Incorrect
  - Invalid
  - Etc. etc.  $\otimes \otimes \otimes$

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

17

## Example I

Consider this execution sequence:

Thread #1	Thread #2
<b>varriable++;</b>	<b>varriable--;</b>
LDA #0x0111	LDA #0x0111
INCA	DECA
STA #0x0111	STA #0x0111

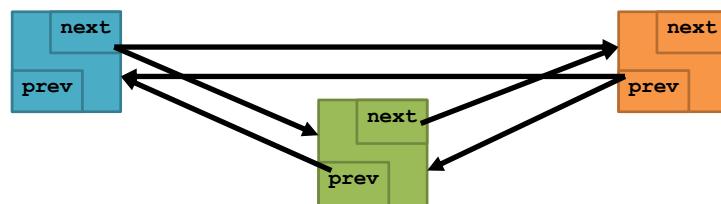
UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

18

18

## Example II

Inserting into a doubly-linked list



**List is inconsistent!**

If reading as a handout, no this list is not inconsistent.  
You'll have to watch the class recording to see the animation.

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

19

## Critical Sections

- Sections of the program (short sequences of ***machine instructions***) that should only be executed by one thread at a time
- Consider a **stack** shared by many threads:
  - All threads should see a consistent state
  - `pop()` and `push()` operations are critical sections
- ***We (the OS) must control execution in all critical sections***

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

20

## The Critical Section Problem

- We need a formalism that allows threads to arbitrate access to critical sections:
  - **Entry section** (used to negotiate entry to the critical section)
  - **Critical section** (do critical task)
  - **Exit section** (signal critical task completed)

## Critical Section Requirements

- **Mutual Exclusion:**
  - Only one thread in critical section at a time
- **Progress:**
  - If a thread wishes to enter an idle critical section, then only threads in entry or exit section may help decide if it is allowed to do so
    - e.g. if a thread doesn't want access to a CS at the moment, it shouldn't stop other threads from accessing it
- **Bounded Wait:**
  - There exists a finite bound on the number of times a waiting thread must "stand-aside"

21

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

22

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

## Software Solutions

- Hmm, tricky...
- Even `variable++` does not translate into an atomic operation
- Typically a **busy-wait** solution is used

## Two Thread Solution I

- Consider a solution involving only two threads (`x` & `y`)
- A first approach is to alternate each thread: `x, y, x, y, x, y, ...`
- **Is this ok?**
  - Which requirements are not met?
  - No **progress** made if thread `x` wants to enter the critical section twice and `y` does not want to enter at all

Shared variables:

`int turn;`

Thread x:

```
while (turn != x) {
    yield();
}
// Critical Section
...
turn = y;
// Non-critical
...
```

23

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

24

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

## Two Thread Solution II

- Modify by adding more state
- Record which threads want to enter their critical section
- Is this ok?
- **NO!** both threads can set their flags to true *"at the same time"* as the assignment is not protected ☺

**Deadlock!**

Shared variables:  
`bool flags[<some size>];`

Thread x:  
`flags[x] = true; ←`  
`while (flags[y]) {`  
 `yield();`  
`}`  
`// Critical Section`  
`...`  
`flags[x] = false;`  
`// Non-critical`  
`...`

## Two Thread Correct Solution

- If solutions I and II are combined, then we can satisfy all three requirements
- Use the boolean array from solution II to record which threads ( $x \& y$ ) wish to enter their critical section
- Resolve deadlock with the alternating approach from solution I

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

25

26

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

## Peterson's Two Thread Solution

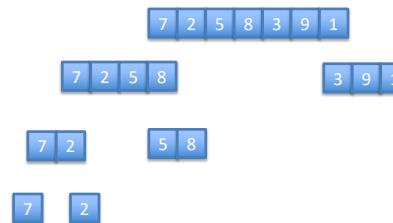
Shared variables:  
`bool flags[<some size>];`  
`int turn;`

Thread x:  
`flags[x] = true;`  
`turn = y;`  
`while (flags[y] && turn == y) {`  
 `yield();`  
`}`  
`// Critical Section`  
`...`  
`flags[x] = false;`  
`// Non-critical`  
`...`

Thread y:  
`flags[y] = true;`  
`turn = x;`  
`while (flags[x] && turn == x) {`  
 `yield();`  
`}`  
`// Critical Section`  
`...`  
`flags[y] = false;`  
`// Non-critical`  
`...`

## Assignment 1

- Merge Sort



<https://opendsa-server.cs.vt.edu/embed/mergesortAV>

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

27

[https://en.wikipedia.org/wiki/Peterson%27s\\_algorithm](https://en.wikipedia.org/wiki/Peterson%27s_algorithm)

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

35

## Assignment 1



```

merge(array, left, middle, right) { /* magic! */ }

mergeSort(array, left, right) {
    if (left < right) {
        middle = (right-left)/2
        mergeSort(array, left, middle)
        mergeSort(array, middle+1, right)
        merge(array, left, middle, right)
    }
}

```

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

49

## Next Lecture

- Introduction to Concurrency Cont.
- Semaphores

UCSC BSOE CMSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

50