

PRINCIPLES OF COMPUTER SYSTEMS DESIGN

CSE130

Winter 2020

Concurrency II - Semaphores



Notices

- Assignment 1 due 23:59 **Sunday January 26**
- Lab 2 due 23:59 **Sunday February 9**
- ~~New Section Times~~
 - ~~Mondays 1:00-2:30pm & 6:30-8:00pm~~
 - ~~Tuesdays 5:00-6:30pm & 6:00-7:30pm (has overlap, may be wrong)~~
 - ~~Wednesdays 2:00-3:30pm & 3:30-5:00pm~~
 - ~~No sections on Thursdays & Fridays~~
- Quiz - Possible Midterm Cancellation
 - On Canvas now
 - Take it, have your say!

2

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Today's Lecture

- Recap
- Hardware Support
- Concurrency Primitives
- Semaphores

- Assignment 1 Secret Sauce
- A little Lab 2 Secret Sauce



Concurrency in Operating Systems

- **"The actual or apparent simultaneous execution of threads"**
 - Actual = multicore and/or multiprocessors
 - Apparent = time slicing
 - Today's Reality = both
- Threads running concurrently in the same process typically want to access shared data, but sharing data in an uncontrolled fashion allows that shared data to become:
 - Inconsistent
 - Incorrect
 - Invalid
 - Generally messed up ☺

4

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

3

1

Critical Sections

- Formalism:
 - **Entry Section**
 - **Critical Section**
 - **Exit Section**
- Requirements:
 - **Mutual Exclusion**
 - **Progress**
 - **Bounded Wait**
- Advice:
 - Read the textbook and review previous lecture handouts for details
 - This is crucial, fundamental information; you MUST know and understand it

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

5

Hardware Assistance

- **Problem:** Software solutions are:
 - Complicated
 - NOT atomic
 - NOT generalizable to n-threads
- **Solution:** Provide synchronization primitives
 - Machine instructions (hardware/microcode)
 - System calls (aid programmers)
 - The OS needs to help, as SW at user level is not enough!

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

6

Test and Set

- Sets boolean lock variable
- Returns original value
- Implemented in hardware
- Is atomic
- **C equivalent:** (not actually implemented in C, it wouldn't be atomic if it was)

```
bool testAndSet(bool *lock) {
    bool tmp = *lock;
    *lock = true;
    return tmp;
}
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

7

Using Test and Set

Shared variables:
`bool lock = false;`

Thread n:
`while (testAndSet(&lock)) {
 yield();
}`
`// Critical Section`
`...`
`lock = false;`
`// Non-critical`
`...`

```
bool testAndSet(bool *lock) {
    bool tmp = *lock;
    *lock = true;
    return tmp;
}
```

lock	Set	Action
F	T	Proceed
T	T	Wait

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

8

2

Swap

- Exchanges two values
 - One private, one shared
- Implemented in hardware
- Is atomic
- C *equivalent*: (again, not actually implemented in C)

```
void swap(bool *a, bool *b) {
    bool tmp = *a;
    *a = *b;
    *b = tmp;
}
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Using Swap

```
void swap(bool *a, bool *b) {
    bool tmp = *a;
    *a = *b;
    *b = tmp
}
```

Shared variables:
`bool lock = false;`

Thread n:

```
bool key = true;
while (key) {
    swap(&lock, &key)
}
// Critical Section
...
lock = false;
// Non-critical
...
```

}

key	lock	key'	lock'	
T	F	F	T	Proceed
T	T	T	T	Wait

9

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

10

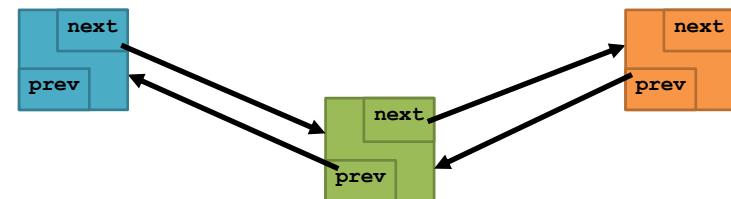
Meet the Requirements?

- **Test And Set & Swap** both satisfy
 - Mutual exclusion
 - Progress
- **Bounded Wait** is not satisfied
 - But with some effort we can add additional structures to ensure this
- Unfortunately...
 - Chip manufacturers don't always include such instructions ☺
- **This lecture we'll look at ways the OS supports safe concurrent execution amongst threads**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Concurrency Primitives

- **Concurrency Primitives**: Mechanisms provided by the OS to allow threads to safely work concurrently
- Inserting into a doubly-linked list is a **classic example** of the need for concurrency primitives



11

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

12

Semaphores

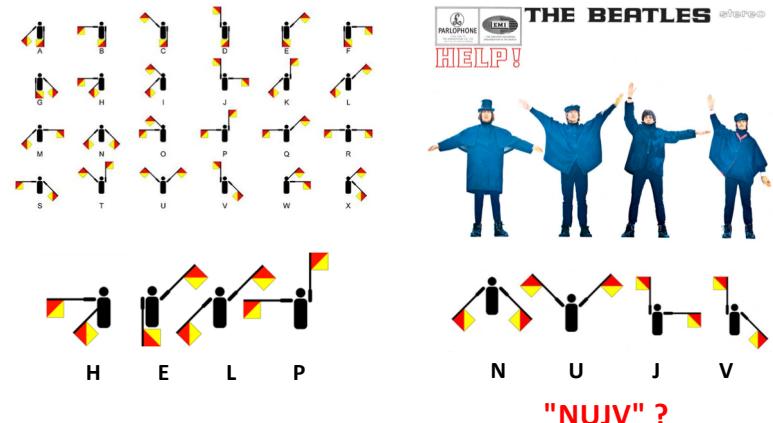
- Dijkstra* recognized the need to manage concurrency amongst threads and processes in 1962**
- But hardware support was far from universal
- In place of machine instructions he proposed the **semaphore**

* Edsger W. Dijkstra 1930 - 2002 : Dutch computer scientist and an early pioneer in many research areas of computing science.
 ** <http://www.cs.utexas.edu/users/EWD/EWD00xx/EWD35.PDF>

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

13

Flag Semaphore



14

Semaphores (Dijkstra -1965)

- P() or wait() "Proberen" (test) a.k.a. "down"
- V() or signal() "Verhogen" (increment) a.k.a. "up"
- P and V are both atomic
- Pseudo code:

```
P(int mutex) {           V(int mutex) {
    while(mutex <= 0);      mutex++;
    mutex--;
}
```

- What's a mutex? MUTual EXclusion

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

15

Using Semaphores

Shared variables:
`int mutex = 1;`

Thread n:

```
...
...           // Non-critical
P(mutex);   // Entry Section
...
...           // Critical Section starts
...
...           // Critical Section ends
V(mutex);   // Exit Section
...
...           // Non-critical again
...
```

Remember:
 "down" "up"
`P(int mutex) {` `V(int mutex) {`
 `while(mutex <= 0);` `mutex++;`
 `mutex--;` `}`

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

16

Counting Semaphores

- The semaphore can also be used to allow N threads into a critical section
- Initialize `mutex` to N
- First N threads decrement `mutex` until `mutex == 0`
- Additional threads ($> N$) must wait
- Ideal for controlling access to a fixed number of resources

Binary Semaphores

- A binary semaphore (a “lock”) limits mutex values to 0 or 1
- Enforces exclusive entry into critical sections
- **Problem:** we are at the mercy of balanced `P()` and `V()` calls
 - i.e. sloppy programmers can mess things up ☺

17

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

18

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Semaphore Limitations

- Semaphores are useful and general tools
- But they do have drawbacks:
 - Programmers must use them correctly
 - Cannot test busy without blocking
 - We can't simultaneously wait on any of several semaphores
 - Indefinite blocking
 - Cooperative only

Producer / Consumer

- Buffering a potentially infinite amount of data into limited storage (memory) is a standard programming problem
- **Producers** (of the data)
 - Devices, networks, processes
- **Consumers** (of the data)
 - User processes / threads
- Two models:
 - Synchronous
 - Direct “hand over” of data
 - Asynchronous
 - Buffered

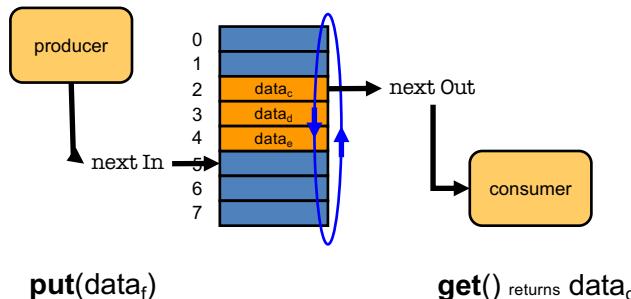
19

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

20

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Circular Bounded Buffers



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

21

Semaphore Synchronisation

Consider a solution to the **producer-consumer** problem using counting semaphores to allocate slots in a circular buffer of integers of size N ...

```
Binary Semaphores: mutex = 1;
Counting Semaphores: vacant = N, occupied = 0; // no. of slots in each state
Buffer Indexes: nextIn = 0, nextOut = 0;
```

```
void put(int x) {
    P(vacant); // are there some vacant slots?
    P(mutex); // am I allowed to change the buffer?
    buffer[nextIn] = x;
    nextIn = (nextIn++) % N;
    V(mutex); // allow others to change the buffer
    V(occupied); // increment the no. of used slots
}
int get() {
    P(occupied); // is there at least one piece of data?
    P(mutex); // am I allowed to change the buffer?
    int x = buffer[nextOut];
    nextOut = (nextOut++) % N;
    V(mutex); // allow others to change the buffer
    V(vacant); // increment the no. of available slots
    return x;
}
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

22

POSIX Semaphores

```
sem_init Initialise an unnamed semaphore
sem_open Initialise and open a named semaphore
sem_wait Lock a semaphore ( put it down, Dijkstra's P )
sem_post Unlock a semaphore ( put it up, Dijkstra's V )
sem_getvalue Get the value of a semaphore
sem_close Close a named semaphore
sem_destroy Destroy an unnamed semaphore
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

23

```
/*
 * buffer-semaphore.c
 *
 * Circular Bounded Buffer protected by unnamed POSIX semaphores.
 *
 * Copyright (C) 2015-2020 David C. Harrison. All rights reserved.
 *
 * You may not use, distribute, publish, or modify this code without the
 * express written permission of the copyright holder.
 *
 * Producers are fast, consumers are slow. More producers than consumers
 * requires producers to wait (block) for slots to become available.
 *
 * To compile:
 *   gcc -o buffer-semaphore buffer-semaphore.c -lpthread -Wall
 *
 * To run:
 *   ./buffer-semaphore
 *
 * Only works on UNIX and UNIX-like systems, so Linux and macOS are fine.
 *
 * If you're on Windows, you have my condolences. \_(ツ)_/-
 */

```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

24

```

#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <semaphore.h>
#include <pthread.h>
#include <unistd.h>

#define PRODUCERS 2
#define CONSUMERS 5

#define MAX_SLEEP_MILLISECONDS 3000
#define MAX_PUTS_PER_PRODUCER 1000

#define BUFFER_SIZE 10

typedef struct buffer_t {
    int buf[BUFFER_SIZE]; /* shared buffer */
    int next_in; /* next slot to add an element to (may not currently be vacant) */
    int next_out; /* next slot to get an element from (may not currently be occupied) */
    sem_t occupied; /* counting semaphore for occupied slots */
    sem_t vacant; /* counting semaphore for vacant slots */
    sem_t mutex; /* binary semaphore to facilitate mutually exclusive access to the buffer */
} Buffer;

Buffer buffer;

static void sleep_ms(int msec) {
    struct timespec ts;
    ts.tv_sec = msec / 1000;
    ts.tv_nsec = (msec % 1000) * 1000000;
    nanosleep(&ts, NULL);
}

/* Initialise the circular bounded buffer BUFFER. */
void buffer_init(Buffer *buffer) {
    sem_init(&buffer->occupied, 0, 0);
    sem_init(&buffer->vacant, 0, BUFFER_SIZE);
    sem_init(&buffer->mutex, 0, 1);
}

/* Blocking insert of VALUE into BUFFER. */
void buffer_put(Buffer *buffer, int value) {
    sem_wait(buffer->vacant); // wait until at least one vacant slot
    sem_wait(buffer->mutex); // wait for exclusive access to the shared buffer

    buffer->buf[buffer->next_in] = value;
    buffer->next_in = (buffer->next_in+1) % BUFFER_SIZE;

    sem_post(buffer->mutex); // release shared buffer so other threads can use it
    sem_post(buffer->occupied); // decrement the number of occupied slots
}

/* Blocking retrieval of next BUFFER entry */
int buffer_get(Buffer *buffer) {
    sem_wait(buffer->occupied); // wait until at least one slot has a valid entry
    sem_wait(buffer->mutex); // wait for exclusive access to the shared buffer

    int value = buffer->buf[buffer->next_out];
    buffer->next_out = (buffer->next_out+1) % BUFFER_SIZE;

    sem_post(buffer->mutex); // release the shared buffer so other threads can use it
    sem_post(buffer->vacant); // decrement the number of vacant slots

    return value;
}

```

25

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

26

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

```

static void *produce(void *arg) {
    long tid = (long) arg;
    int value;
    printf("P %ld running\n", tid);

    for (int i = 0; i < MAX_PUTS_PER_PRODUCER; i++) {
        value = tid*1000;
        printf("=> P %ld produced %04d\n", tid, value);
        buffer_put(&buffer, value);
    }

    printf("P %ld ### finished ##\n", tid);
    pthread_exit(NULL);
    return NULL;
}

static void *consume(void *arg) {
    long tid = (long) arg;
    int value, millis;
    printf("C %ld running\n", tid);

    for (;;) {
        millis = rand() % MAX_SLEEP_MILLISECONDS;
        value = buffer_get(&buffer);
        printf("<- C %ld consumed %04d - sleeping for %dns\n", tid, value, millis);
        sleep_ms(millis);
    }

    // Will never return, consumers are in an infinite loop
    return NULL;
}

```

27

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

28

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

29

Self Study (totally not graded, just do it for fun!)

- Copy the source code from handouts
 - Build it, run it, study it
 - Change it so producers are slow and consumers are fast
 - **Q: What happens?**
 - Increase no. of producers & decrease no. of consumers
 - **Q: What happens?**
 - Take it to parties in an attempt to impress your friends
 - **Q: What happens?**
 - **A: You no longer have any friends.**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

30

Assignment 1 - Secret Sauce

- Using POSIX Shared Memory System Calls
 - Re-install and add -lrt flag to Makefile.libs
 - When a process calls `fork()`
 - How similar are the parent and child processes?
 - Do they share a program counter?
 - When does the program counter of each change?
 - Do they share executable code?
 - What happens when they each arrive at a return statement?
 - Watch your processes!!
 - Two users where chewing up 64,000 processes between them ☹
 - See how many processes you have and kill 'em:

```
$ ps aux | grep <cruzid>
$ killall -u <cruzid>
```



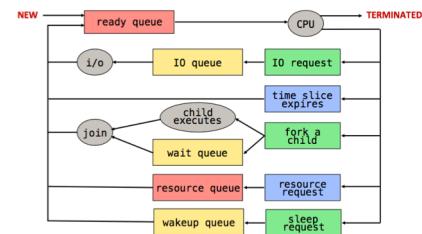
UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Lab 2 - Secret Sauce

- We have not yet covered in lectures everything you need to understand, but there's a whole bunch of stuff we do know:
 - Pintos is a time-sharing OS
 - When a **time-slice** expires:
 - The running (current) thread is taken off the CPU and placed at the back of the **ready queue**
 - The thread at the front of the ready queue is then removed and placed on the CPU

The diagram illustrates the state transitions of threads in Pintos. It shows a vertical stack of four main states: NEW, ready queue, IO queue, and CPU. Transitions are indicated by arrows:

 - From NEW to ready queue.
 - From ready queue to IO queue via an i/o transition.
 - From ready queue to CPU via a time slice expires transition.
 - From IO queue to CPU via an IO request transition.
 - From CPU to ready queue via a time slice expires transition.
 - From CPU to child executes via a fork child transition.
 - From child executes to wait queue via a join transition.
 - From wait queue to CPU via a resource request transition.



- If we want to implement **priority based scheduling** we have to sort the ready queue as and when appropriate

UCSC BSOE CSE120 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

32

Next Lecture

- Condition Variables

33

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.