

PRINCIPLES OF COMPUTER SYSTEMS DESIGN

CSE130

Winter 2020

Concurrency IV - Deadlocks



Notices

- Assignment 1 due 23:59 **Monday January 27 (24 Hour Extension)**
- Assignment 2 due 23:59 **Sunday February 2**
- Lab 2 due 23:59 **Sunday February 9**
- Section Times Reminder
 - Monday 1-2:30pm & 6:30-8pm
 - Tuesday 5-6:30pm
 - Wednesday 2-3:30pm
 - Friday 11am-12:30pm & 12:30-2pm
- Quiz - Possible Midterm Cancellation
 - On Canvas now
 - Take it, have your say!

2

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Today's Lecture

- Introduction to Deadlocks
- The Dining Philosophers Problem
- Race Conditions
- Resource Allocation Graphs
- Deadlock Prevention, Avoidance, and Detection
- Introduction to Assignment 2

Deadlock Defined

"A situation from which it is impossible to proceed"

- Operating Systems must be designed to deal with deadlocks as resources are finite
- Not all OSs deal with deadlocks in the same way
- Best known illustration is “The Dining Philosophers Problem”

4

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

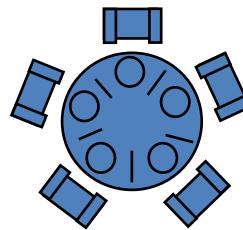
3

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

1

The Dining Philosophers Problem

- Five dining philosophers
- Five chopsticks
- Philosophers are either thinking, or eating
- To eat, each philosopher needs two chopsticks
- Chopsticks are put down while thinking



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

5

Codifying The Dining Philosophers Problem

- A simple solution is to represent each chopstick by a binary semaphore
- Then use an array of semaphores indexed by position
- A chopstick is picked up with:
P(chopsticks[pos])
- And set down with:
V(chopsticks[pos])

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

6

Simple Solution - Pseudo Code

```
while (true) {
    P(chopsticks[left]);
    P(chopsticks[right]);
    // nom, nom, nom
    V(chopsticks[left]);
    V(chopsticks[right]);
    // ponder life, the universe, and everything
}
```

Problem?

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

7

Deadlocked Philosophers

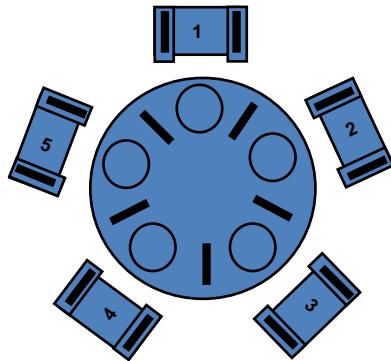
- What if the philosophers all pick up their left chopstick at the same time?
 - No philosopher can grab their right chopstick (another's left chopstick) and will therefore wait forever and starve
- This is **the classic example of deadlock**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

8

2

Not Deadlocked Philosophers



(see lecture video for animation)

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

9

Deadlocked Philosophers



(see lecture video for animation)

10

Race Conditions

- The dining philosophers are **in a race** to get the chopsticks
- Race Condition:**
"A flaw in a system or process whereby the output of the process is unexpectedly and critically dependent on the sequence or timing of other events."
- In the Pintos source code you'll see at least one comment observing that some function call sequences may be "racy"

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

11

Deadlock Prevention

- How can we stop the dining philosophers from becoming deadlocked and starving to death?
- Possibilities:**
 - Limit the number of philosophers to 4 or 5 at the table at any one time
 - Pick up the chopsticks only if both are available
 - i.e. define a critical section
- Best Solution:**
 - Alternate the order that chopsticks are picked up if the philosopher is seated at an odd or even seat

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

12

Self Study (totally not graded, but fun!)

- Code up The Dining Philosophers Problem using POSIX Semaphores
- See if you can get your implementation to deadlock
 - How frequently does it deadlock?
- Try the solutions from the previous page
 - Did they fix the problem?
 - How do you know they've fixed the problem?

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

13

Operating System Definitions

- **Deadlock:**
"A set of processes is in a deadlock state when every thread or process in the set is waiting for an event that can only be caused by another thread or process in the set."
 - Threads competing for **critical sections**
 - Processes or threads competing for **resources**
 - Resource (CPU, memory, tape drives, printers, network cards...)
 - Some resources are **preemptable**, some are not
 - Resource access lifecycle:
 - **request** (wait), **use**, **release** (signal)
- **Starvation:**
 - A thread or process waiting forever

Deadlock => Starvation
Starvation \neq Deadlock

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

14

Deadlock?

MUTEX x = 1, y = 1; // shared between threads

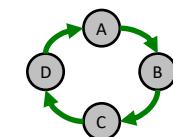
Thread 1	Thread 2
P(x);	P(y);
P(y);	P(x);

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

15

Necessary Conditions for Deadlock

- **Mutual exclusion**
 - At least one resource does not support simultaneous use
- **No preemption**
 - A resource cannot be seized from the holder
- **Hold and wait**
 - At least one process or thread must hold a resource, and be waiting on another
- **Circular wait**
 - e.g. A waits on B waits on C waits on D waits on A

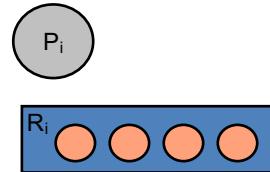


UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

16

Resource Allocation Graphs

- Useful in examining deadlocks
- Set of processes $[P_1, P_2, \dots]$
- Set of resources $[R_1, R_2, \dots]$
 - Can be more than one instance of some resources
 - Resources are identical and interchangeable



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

17

Resource Allocation Graphs

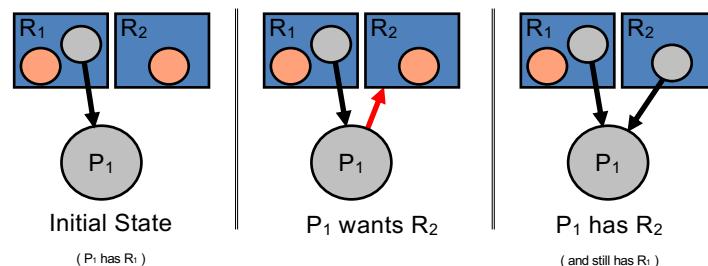
- Protocol:
1. Request a resource
 2. Acquire an instance of a resource (may have to wait to get it)
 3. Use the resource
 4. Release the resource
- Resource may be a lock, a buffer, a device, etc.
 - “Acquire” and “Release” are system calls
 - Processes may need to wait on a resource
 - Typically processes wait for an event associated with the release of the resource

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

18

Constructing Resource Allocation Graphs

- Draw:
 - Edges from P_i to R_j on **request** (P_i, R_j) “wants” →
 - Edges from R_j to P_i on **allocation** (R_j, P_i) “has” →



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

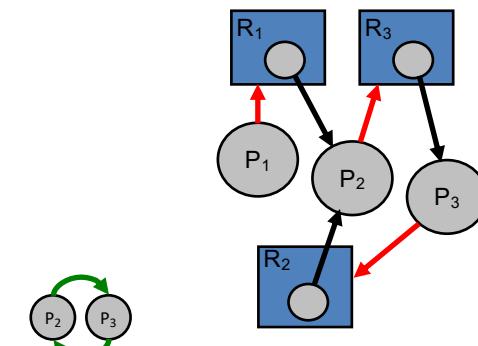
19

Deadlock #1

If:
 P_1 wants R_1
 P_2 wants R_3
 P_3 wants R_2

And:
 P_2 has R_1
 P_2 has R_2
 P_3 has R_3

Then:
 P_1 waiting for R_1
 P_2 waiting for R_3
 P_3 waiting for R_2

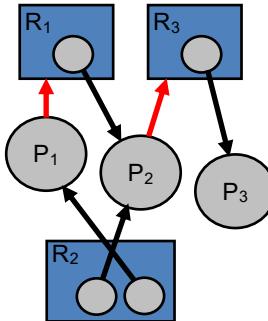


Deadlock #2

P₁ waiting for R₁
P₂ waiting for R₃
P₃ not waiting for anything
=> P₃ runnable

Deadlock? **NO!**

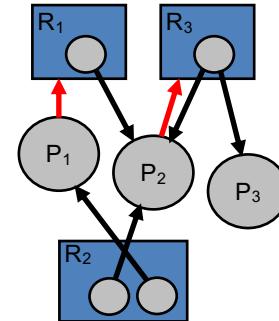
But there may be a deadlock a little later...
...it depends of what happens next



Deadlock #2 - Possibility 1

P₃ releases R₃
P₂ acquires R₃
=> P₂ runnable

Deadlock? **NO!**



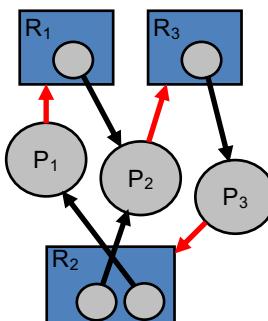
(see lecture video for animation)

Deadlock #2 - Possibility 2

P₃ requests R₂

Deadlock? **YES!**

Circular Waits between P₁, P₂, P₃ and P₃, P₂
even though two instances of R₂



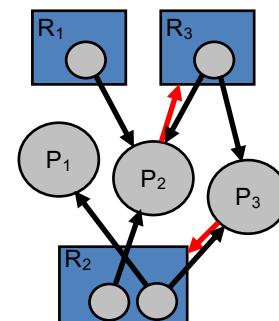
(see lecture video for animation)

Is Circular Wait Sufficient for Deadlock?

P₂ and P₃ are in a circular wait
P₁ releases R₂
=> P₂ and P₃ no longer in a circular wait
P₃ acquires R₂
=> P₃ runnable

Deadlock? **NO!**

P₃ releases R₂ and R₃
P₂ acquires R₃
=> P₂ runnable



(see lecture video for animation)

Handling Deadlocks

- **Ignore** them (deadlocks are rare)
 - Cheap, can rely on manual detection
 - Avoids performance hits
- Make sure they don't happen:
 - **Prevention:** Limit ways to request resources
 - **Avoidance:** Allocate resources to stop them happening
- **Detect** them and recover

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

25

Deadlock Prevention

- Eliminate **mutual exclusion**:
 - Not generally possible
- Eliminate **no preemption**:
 - If we wait on a resource held by another waiting process, then preempt the resource and allocate to us
- Eliminate **hold and wait**:
 - If waiting on a resource, release all currently held resources
 - Resource utilisation may be poor
 - Starvation may occur with popular resources
- Eliminate **circular wait**:
 - Impose progressive ordering for resources requests R_1, R_2, \dots, R_n
 - Multiple resources of the same type must be requested together

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

26

Deadlock Prevention

- Difficult to achieve without it leading to low resource utilization which leads to low throughput, i.e. poor performance
- Operating Systems generally do not try to prevent deadlocks

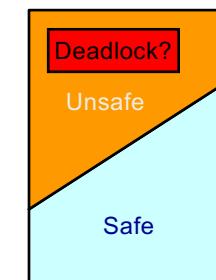
UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

27

Deadlock Avoidance

"A resource allocation is safe if (and only if) there exists an execution sequence $\langle P_a, P_b, \dots, P_n \rangle$ such that all processes can run to completion"

- Takes into account resources the process might use
- All other states are unsafe
- From a safe state the system can allocate resources to avoid deadlock
- Requires a priori knowledge of maximum resources required by each process



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

28

Deadlock Avoidance - Example

12 instances of a single resource are available:

max = most instances of the resource a process P_n might ask for
alloc = instances of the resource currently allocated to P_n
free = instances not currently allocated to any process
risk = instances P_n could ask for in next timeslot

	max	alloc	risk
P_0	10	5	5
P_1	4	2	2
P_2	9	2	7
free	3		

$\langle P_1, P_0, P_2 \rangle$ is **safe** ☺

When P_1 finishes, 5 free, P_0 can run

When P_0 finishes, 10 free, P_2 can run

	max	alloc	risk
P_0	10	5	5
P_1	4	2	2
P_2	9	3	6
free	2		

All sequences are **unsafe** ☹

Deadlock if P_0 requests 5

Deadlock if P_2 requests 6

(only 4 free when P_1 finishes)

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

29

Making Sure Deadlocks Don't Happen

- “Unsafe” does not mean a deadlock *will* occur, rather there is a *possibility* of deadlock
- It’s about guaranteeing deadlocks don’t happen:
 - State max resource requirements up-front
 - Don’t grant R_i to P_j unless the resulting state remains safe
 - Can allocate max if needed
- However, resource utilization may be negatively impacted
 - There’s always an overhead to safety checking
- Various approaches:
 - Resource Allocation Graphs
 - Dijkstra’s Bankers’ Algorithm (self study, Google it!)

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

30

Deadlock Detection

- If we don’t want to, or can’t avoid / prevent deadlocks, we can detect them and do something about it:
- Break the circular wait:**
 - Option 1: Kill all deadlocked processes/threads
 - Option 2: Kill deadlocked processes/threads one-by-one
 - After each kill, see if the deadlock is broken
 - Remember, some processes should not be killed!
- Preempt resources:**
 - Rollback to a checkpoint
 - Run the process in a different sequence

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

31

Self Study (who am I kidding, no one is actually going to do this)

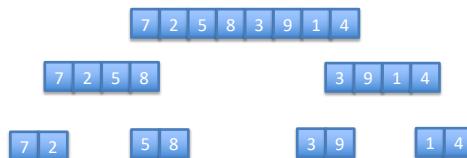
- Modify your code for the Dining Philosophers Problem to detect deadlock and either **break the circular wait** or **preempt** the chopsticks and rollback

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

32

Assignment 2

- Multi Threaded Merge Sort



<https://opendsa-server.cs.vt.edu/embed/mergesortAV>

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

33

POSIX Thread Functions (not a complete list)

Process "equivalent"		
pthread_create	Create a new thread	fork
pthread_join	Join with terminated thread	wait
pthread_exit	Terminate calling thread	exit

http://man7.org/linux/man-pages/man2/shmget_2.html

34

Next Lecture

- CPU Scheduling I

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

35