

PRINCIPLES OF COMPUTER SYSTEMS DESIGN

CSE130

Winter 2020

CPU Scheduling I



Notices

- Assignment 2 due 23:59 **Sunday February 2**
- Lab 2 due 23:59 **Sunday February 9**

2

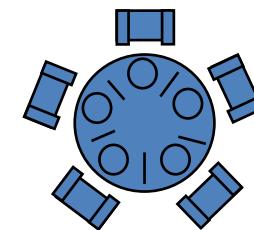
Today's Lecture

- The Dining Philosophers Problem Revisited
- Introduction to Scheduling
- Scheduling Opportunities
- Scheduling vs. Dispatching
- First Come First Served
- Round Robin
- Assignment 2 Secret Sauce



The Dining Philosophers Problem

- Five dining philosophers
- Five chopsticks
- Philosophers are either thinking, or eating
- To eat, each philosopher needs two chopsticks
- Chopsticks are put down while thinking



4

Codifying The Dining Philosophers Problem

- A simple solution is to represent each chopstick by a binary semaphore
- Then use an array of semaphores indexed by position
- A chopstick is picked up with:
 $\text{P}(\text{chopsticks[pos]})$
- And set down with:
 $\text{V}(\text{chopsticks[pos]})$

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

5

Simple Solution - Pseudo Code

```

while (true) {
    P(chopsticks[left]);
    P(chopsticks[right]);
    // nom, nom, nom
    V(chopsticks[left]);
    V(chopsticks[right]);
    // ponder life, the universe, and everything
}

```

If all philosophers pick up their left chopstick at the same time, we get DEADLOCK as none of them can now pick up their right chopstick ☺

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

6

```

#define PHILOSOPHERS 5
#define MIN_EAT_MS 500
#define MAX_EAT_MS 1500
#define MIN_THINK_MS 1000
#define MAX_THINK_MS 3000

// how fast a philosopher can pick up a chopstick,
// / 0 = instantaneous
static int muscle_response_time = 0; ←

typedef struct philosopher_t
{
    int pid; // philosopher (not process) id
    int nomes; // how many times philosopher has eaten
    sem_t *left; // binary semaphore for left chopstick
    sem_t *right; // binary semaphore for right chopstick
} philosopher_t;

static sem_t *chopsticks[PHILOSOPHERS];
static philosopher_t philosophers[PHILOSOPHERS];

static void sleep_ms(int msec) {
    if (msec > 0) {
        struct timespec ts;
        ts.tv_sec = msec / 1000;
        ts.tv_nsec = (msec % 1000) * 1000000;
        nanosleep(&ts, NULL);
    }
}

static void monitor() {
    for (;;) {
        sleep_ms(1000);
        printf("\n");
        for (int i = 0; i < PHILOSOPHERS; i++) {
            printf("Philosopher %d has eaten %d times\n",
                  philosophers[i].pid, philosophers[i].nomes);
        }
    }
}

static void philosopher_main(void *arg)
{
    philosopher_t *philosopher = (philosopher_t*) arg;
    for (;;) {
        sem_wait(philosopher->left);
        sleep_ms(muscle_response_time); ←
        sem_wait(philosopher->right);

        sleep_ms(MIN_EAT_MS + rand() % (MAX_EAT_MS - MIN_EAT_MS));

        sem_post(philosopher->left);
        sem_post(philosopher->right);

        philosopher->nomes++;
        sleep_ms(MIN_THINK_MS + rand() % (MAX_THINK_MS - MIN_THINK_MS));
    }
}

int main(int argc, char* argv[])
{
    muscle_response_time = argc > 1 ? atoi(argv[1]) : 0;
    srand(getpid());
    for (int i = 0; i < PHILOSOPHERS; i++)
        sem_init(&chopsticks[i], 0, 1);

    pthread_t junk;
    for (int i = 0; i < PHILOSOPHERS; i++) {
        philosopher_t pid = i;
        philosopher_t.right = &chopsticks[i];
        philosopher_t.left = &chopsticks[i+1] % PHILOSOPHERS;
        pthread_create(
            &junk, NULL, philosopher_main,
            (void*) &philosophers[i]);
        monitor();
    }
}

```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

7

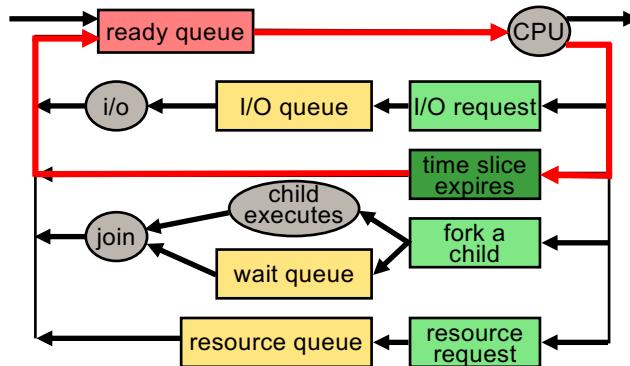
Introduction to CPU Scheduling

- Efficient use of the CPU is critical to overall performance of the entire system
- The CPU is fundamentally a pre-emptible resource
- **Arguably the most important role of the operating system is arbitrating access to the CPU**
- Metrics including priority, previous activity, etc. can be used in determining which process next gets use of the CPU

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

8

Preemptive Scheduling (no IO wait)



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

9

Scheduling Opportunities

- **Running thread yields** or:
 - Makes a system call: `sleep()`, `fork()` etc.
 - Has to wait for a semaphore, lock, mutex, network socket, etc.
- **Thread terminates**
- **An interrupt occurs**
 - Current thread is still ready
 - Blocked thread(s) become(s) ready
- **Preemption**
 - Usually based on timer interrupt
 - Choose from all ready threads (how you choose is the *scheduling algorithm*)
- **Non-preemption**
 - Continue with current thread until it finishes ☺

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

10

Scheduling vs. Dispatching

- **A Scheduler:**
 - Maintains the ready queue
 - Selects next process to run from those in the ready queue
- **A Dispatcher:**
 - Puts the selected process onto the CPU by:
 - Switching context (i.e. reloading the PCB)
 - Switching to user mode (i.e jumping to the proper location in the executable code to restart that program where it left off at the previous context switch)
 - **Dispatch Latency:**
 - Time it takes for the dispatcher to stop one process and start another running - an important OS metric
- **NOTE:** The terms “Scheduler” and “Dispatcher” are frequently used interchangeably and inconsistently – try to avoid doing this

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

11

Scheduling Criteria

- **Utilization**
 - Try to keep the CPU as busy as possible (used / available)
- **Throughput**
 - # of threads that complete their execution per time unit (threads / time)
- **Turnaround time**
 - Amount of time to execute a particular task start to finish ($t_{\text{finish}} - t_{\text{start}}$)
- **Waiting time**
 - Amount of time a thread has been waiting in the ready queue (t_{readyQ})
- **Response time**
 - Amount of time from when a request was submitted until the first response is produced ($t_{\text{firstresponse}} - t_{\text{start}}$)

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

12

Scheduling Considerations

- Different schedulers will give different priorities to different goals in different situations
- Throughput depends on the size and type of the jobs
- **Good metrics?**
 - Batch Systems: **Turnaround time**
 - Multiprogrammed: **Waiting time**
 - Time Sharing / Interactive: **Response time**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

13

Waiting Time (for any scheduling algorithm)

For a thread **T** that gets scheduled **N** times during its lifetime:

arrival = time **T** initially enters the scheduling queue

start_n = time **T** gets the CPU for the **nth** time

finish_n = time **T** yields or is forced to yield the CPU for the **nth** time

$$\text{Waiting Time} = (\text{start}_1 - \text{arrival}) + \sum_{n=2}^N (\text{start}_n - \text{finish}_{n-1})$$

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

14



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

15

First Come First Served (non preemptive)

- Non-preemptive FCFS is simple, but...
 - Long running, CPU bound threads introduce highly variable wait times
 - **burst** : time slots threads will use before yielding or having to wait for IO
 - **t_{arrive}** : time slot in which the thread first entered the ready queue

	burst ₁	t _{arrive}
T ₁	24	2
T ₂	3	0
T ₃	3	1



Waiting times: $T_1 = 6-2 = 4$ $T_2 = 0$ $T_3 = 3-1 = 2$

Mean waiting time: $(4 + 0 + 2) / 3 = 2$

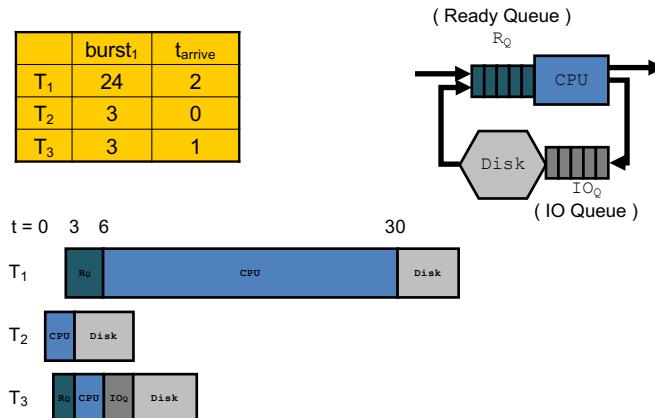
Mean turnaround time: $(3 + (6-1) + (30-2)) / 3 = 12$

Total turnaround time: 30

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

16

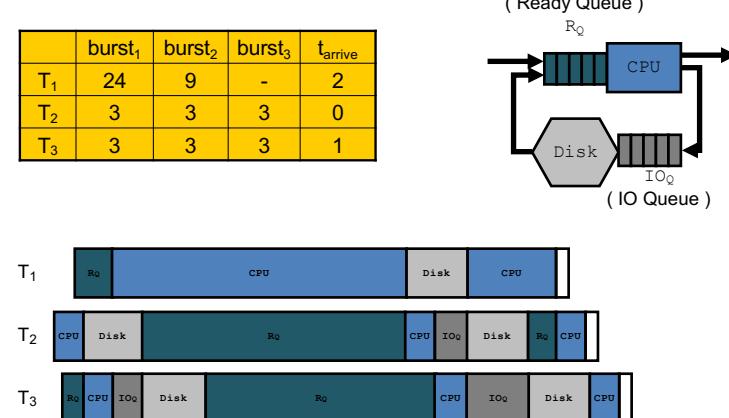
System View of FCFS



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

17

System View of FCFS

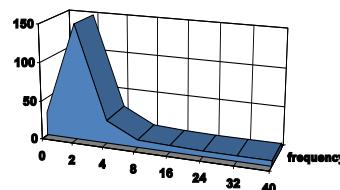


UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

18

Thread Characterisation

- Multiprogramming aims for maximum CPU utilization
- CPU & I/O Burst Cycles
 - Execution consists of cycles of CPU followed by I/O
 - Both bursts may have to wait in a queue
- CPU burst distribution critical to scheduling algorithm



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

19

Round Robin

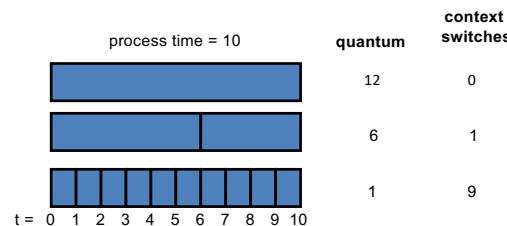
- An attempt to eliminate the convoy effect as they tend to be bad for time-sharing systems
- Each process gets a quantum, q , of CPU time
 - Usually $10 \leq q \leq 100$ milliseconds
- When quantum expires, process is **preempted** & returned to the ready queue
- If there are n processes, each thread gets $\sim 1/n$ of CPU time
- No process waits more than $q*(n-1)$ until its next go on the CPU

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

20

Round Robin Performance

- Care must be taken when choosing q
 - q too large => same as non-preemptive FCFS
 - q too small (\approx time to context switch) => overhead is too high



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

21

Turnaround Time & Quantum I

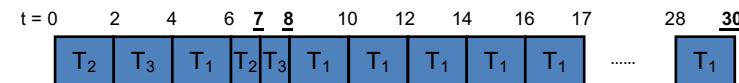
- Turnaround time is impacted by the quantum selected
 - With $q = 2$ (assume context switch is ~instantaneous) what is the mean turnaround time?

	burst ₁	t _{arrive}
T ₁	24	2
T ₂	3	0
T ₃	3	1

Mean turnaround time is:
$$((7-0) + (8-1) + (30-2)) / 3 = 14$$

But that's worse than FCFS ☹

Does it get better if we make $q = 5$?

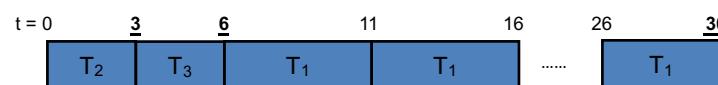


22

Turnaround Time & Quantum II

- Can modifying the quantum improve turnaround time?
 - With $q = 5$ (assume context switch is \sim instantaneous) what is the mean turnaround time?

	burst ₁	t _{arrive}
T ₁	24	2
T ₂	3	0
T ₃	3	1



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

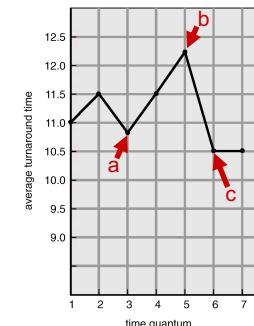
23

Turnaround Time & Quantum III

process	time
P_1	6
P_2	3
P_3	1
P_4	7

If quantum is ten times the length of a context switch, then 10% of CPU time is spent/wasted context switching

- Turnaround time depends on size of the quantum selected
 - Mean turnaround does not always improve with the length of the quantum (**b**)
 - Mean turnaround improved (reduced) if most processes (rule of thumb 80%) finish next CPU burst within the next quantum (**a**) and (**c**)



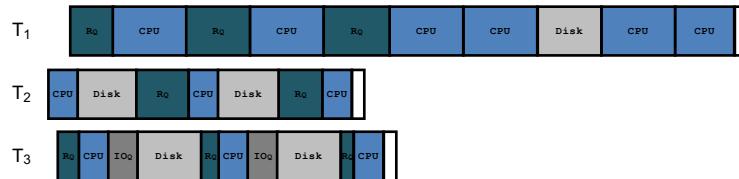
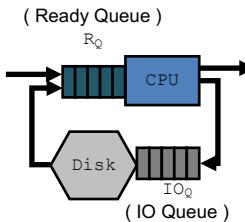
UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

24

System View of RR

	burst ₁	burst ₂	burst ₃	t _{arrive}
T ₁	24	9	-	2
T ₂	3	3	3	0
T ₃	3	3	3	1

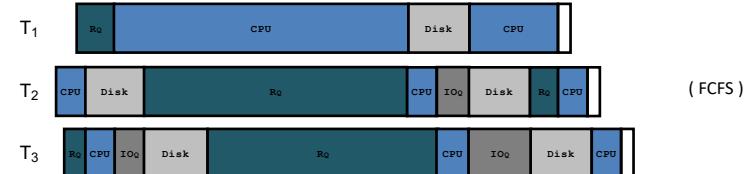
Quantum chosen to exceed most bursts: $q = 5$



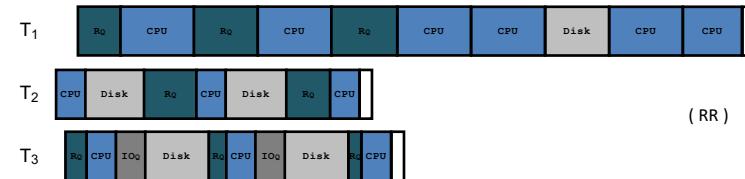
UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

25

Comparison: FCFS v RR



(FCFS)



(RR)

26

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

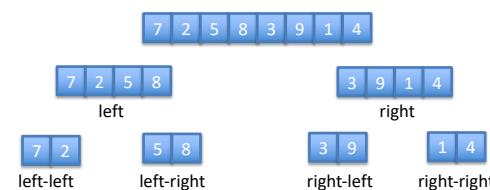
Conclusion?

- FCFS:** Bad guys finish first ☹
 - Threads that hog the CPU tend to prevent others from getting a chance to execute
- RR:** Nice guys finish first ☺
 - Threads that do modest chunks of work before yielding or waiting for IO get more and earlier chances to use the CPU
- In time-sharing situations, RR is usually preferable
- But can we do better than Round Robin?

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

27

Assignment 2 - Secret Sauce



For 2.6 times speed up, how many threads do we need?
What do these threads do?
Where are the merges done?

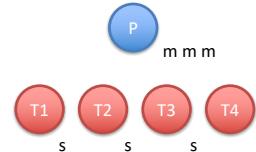
28

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Assignment 2 - Process Pseudo Code 1

```

split array into four
create four threads to sort a 1/4 of the array each
wait for all four threads to finish
merge left-left & left-right
merge right-left & right-right
merge left and right
  
```



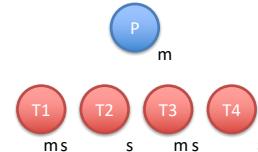
How many merges are running concurrently?



Assignment 2 - Process Pseudo Code 2

```

split array into four
create four threads to sort a 1/4 of the array each
T1 waits for T2 to finish then merges left-left & left-right
T3 waits for T4 to finish then merges right-left & right-right
wait for T1 and T3 to finish
merge left and right
  
```



29

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

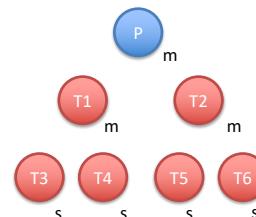
30

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Assignment 2 - Process Pseudo Code 3

```

split array into two
create two threads to sort half of the array each
T1 and T2 each create two threads to sort half of their half
T1 waits for T3 and T4 to finish
  then merges left-left & left-right
T2 waits for T5 and T6 to finish
  then merges right-left & right-right
wait for T1 and T2 to finish
merge left and right
  
```



Next Lecture

- Doing better than Round Robin

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

31

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

32