

PRINCIPLES OF COMPUTER SYSTEMS DESIGN

CSE130

Winter 2020

CPU Scheduling II



Notices

- Assignment 2 due 23:59 Sunday February 2
- Lab 2 due 23:59 Sunday February 9
- No assignment next week, so plenty of time for Lab 2 😊

2

Today's Lecture

- Recap
 - Waiting Time Calculation
 - Simple Scheduling Algorithms: FCFS and RR
- Shortest Job First (non-preemptive & preemptive)
- Advantages and Disadvantages of FCFS, RR, and SJF
- Multilevel Queue
- Multilevel Feedback Queue
- Scheduling in Pintos
- Priority Donation

Waiting Time (for any scheduling algorithm)

For a thread **T** that gets scheduled **N** times during its lifetime:

arrival = time **T** initially enters the scheduling queue

start_n = time **T** gets the CPU for the **nth** time

finish_n = time **T** yields or is forced to yield the CPU for the **nth** time

$$\text{Waiting Time} = (\text{start}_1 - \text{arrival}) + \sum_{n=2}^N (\text{start}_n - \text{finish}_{n-1})$$

Simple Scheduling Algorithms

- **FCFS:** Bad guys finish first ☹
 - Threads that hog the CPU tend to prevent others from getting a chance to execute
- **RR:** Nice guys finish first ☺
 - Threads that do modest chunks of work before yielding or waiting for IO get more and earlier chances to use the CPU
- In time-sharing situations, RR is usually preferable
- But can we do better than Round Robin?

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

5

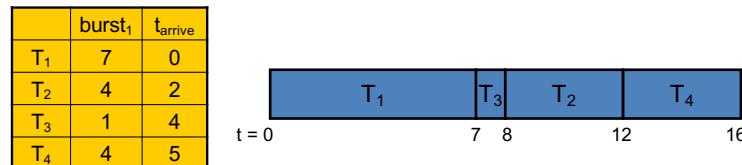
Shortest Job First

- Burst lengths seem to be critical, so...
- Select the ready job (process or thread) with shortest CPU burst
- When a new thread arrives with a shorter burst than the currently running job (the current thread) SJF can be:
 - **Non-preemptive**
 - Don't preempt current thread, let it complete this CPU burst, then run the newly arrived thread, unless... (**unless what?**)
 - **Preemptive**
 - If arriving thread has CPU burst less than remainder of the current job's current burst, preempt (i.e. kick current thread off CPU and run the newly arrived thread)
 - Sometimes called "**Shortest Remaining Time First**" (SRTF)
- **Preemptive SJF** is optimal for time-sharing systems
 - Gives minimum mean waiting time for any given set of threads
 - **What other scheduling metric is likely (but not guaranteed) to be minimised?**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

6

Non-Preemptive SJF



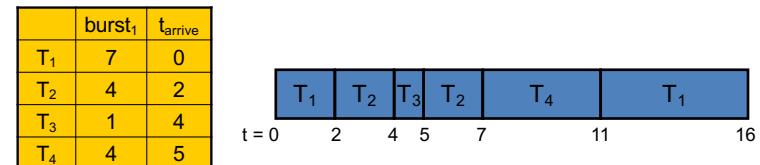
$$\text{Mean waiting time} = ((0-0) + (8-2) + (7-4) + (12-5)) / 4 = 4$$

$$\text{Mean turnaround time} = ((7-0) + (12-2) + (8-4) + (16-5)) / 4 = 8$$

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

7

Preemptive SJF



$$T_1 \text{ waiting time} = (0-0) + (11-2) = 9$$

$$T_2 \text{ waiting time} = (2-2) + (5-4) = 1$$

$$T_3 \text{ waiting time} = 4-4 = 0$$

$$T_4 \text{ waiting time} = 7-5 = 2$$

$$\text{Mean waiting time} = (9 + 1 + 0 + 2) / 4 = 3 \text{ (was 4)}$$

$$\text{Mean turnaround time} = ((16-0) + (7-2) + (5-4) + (11-5)) / 4 = 7 \text{ (was 8)}$$

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

8

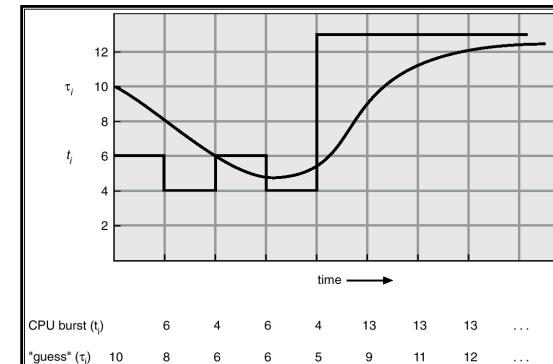
2

Unfortunately...

- SJF needs to see into the future
- Not currently possible ☺
- So we have to make a guess, but how?
- We don't actually know what the next CPU burst will be, so we predict it based on what happened in the past:
 - Decaying weighted mean, $\alpha = \{0 \rightarrow 1\}$
 - $\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau$
 - Exponential mean
 - $\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0$



CPU Burst Prediction



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

9

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

10

Comparing Scheduling Algorithms

- First Come First Served
 - Pros: Simple
 - Cons: Short Jobs can get stuck behind long ones
- Round Robin
 - Pros: Better for short jobs (especially if they are shorter than the quantum)
 - Cons: Not great when jobs are same length (waste time context switching)
- Shortest Job First
 - Pros: Optimal for mean waiting time (i.e good for time-sharing)
 - Cons: Hard to predict the future - unfair on long jobs?

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

11

Multilevel Queue (a.k.a "multiple queues")

- Recognize that threads can be classified into different groups
- The single ready queue becomes, say, two queues:
 - Foreground (interactive)
 - Background (batch)
- Each queue can have its own scheduling algorithm, for example:
 - Foreground - RR
 - Background - FCFS

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

12

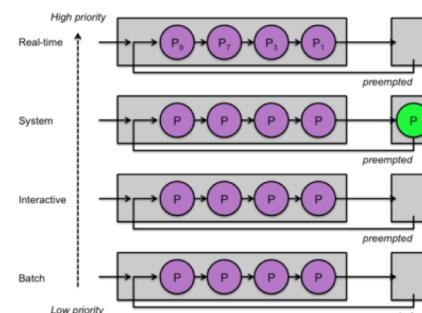
Priority Scheduling

- A **priority** is associated with each thread
 - SJF is a priority scheduler that uses the predicted CPU burst as the priority value
 - Alternatively, we can explicitly assign priorities
- Priority determines which ready queue the thread will join
- CPU is allocated to the thread with the highest priority, as usual, there are two types:
 - **Preemptive**
 - **Non-preemptive**
- Simplest form is **Fixed Priority Scheduling**
 - Serve all from highest priority queue, then all from next highest priority queue, and so on

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

13

Multiple Queues for Multiple Priorities



One problem is that the process needs to be assigned to the most suitable priority queue *a priori*

If a CPU-bound process is assigned to a short-quantum, high-priority queue, that's suboptimal for the process and for overall throughput

Remember: Only one CPU - Diagram is not the best ☺

14

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Multi Level Queues : Problems?

- There is a possibility of starvation
- One solution is to **queue time slice** - each **queue** gets a certain minimum proportion of CPU time which it can schedule amongst its threads, for example:
 - 80% to foreground in RR
 - 20% to background in FCFS
- Another solution is to apply **runtime feedback**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

15

Multilevel Feedback Queue

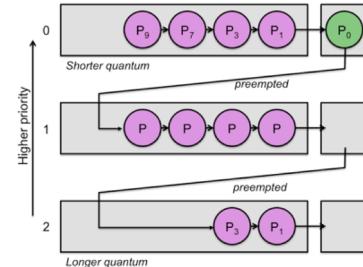
- Introduce scheduling to move jobs between queues
- A **multilevel feedback queue** scheduler is defined by:
 - Number of queues
 - Scheduling algorithms for each of the queues
 - Methods used to determine:
 - **Initial** queue for thread
 - When to **promote** thread
 - When to **demote** thread

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

16

A Simple Queue Demotion Scheme

- Priority is lowered when quantum exceeded
- Q_0 drained before $Q_{>0}$
 - This discriminates against long jobs
 - But improves life for shorter jobs
- Each queue may have its own scheduling algorithm



Remember: Only one CPU - Diagram is not the best ☺

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

17

Queue Demotion : Problems?

- In general, multi-level feedback queues favour I/O bound processes by keeping them at high priorities
- But... what would happen if an I/O intensive process such as a text editor first needed to perform some CPU-intensive work such as initializing a lot of structures and tables?
 - Although it would get scheduled at a high priority initially, it would find itself using up its quantum and demoted to a lower level
 - There too, it may not be done with its tasks and be demoted to yet lower levels before it blocks in I/O
 - Now it will be doomed to run at a low priority level, providing potentially inadequate interactive performance ☺

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

18

Queue Promotion

- Process Aging** can rescue an interactive program that had a period of CPU-intensive activity and was "punished" down to a lower priority queue
 - The unfortunate process is promoted back to a higher priority when it gets older than some threshold
- Alternatively, we can raise the priority of a process whenever it becomes blocked on I/O

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

19

Pintos Scheduling

- Pintos implements a single ready queue with a quantum
- In Lab 2, you are implementing a priority scheduler ☺
- There are 64 priorities, default is 31, 0 is lowest
- Threads can change priorities
- Every queue (ready, and any related to concurrency primitives like semaphores and condition variables) must reflect these priorities when removing (or adding) elements
- General principle is that no higher priority thread should have to wait while a lower priority thread executes**
- Preemption is super important, so get it right!**
- But priority donation is problematic ☺

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

20

Priority Donation : Example Pseudo Code

```
Thread A (priority 20) {
    acquire Lock
    create Thread B
    create Thread C
    release Lock
}

Thread B (priority 40) {
    acquire Lock
    release Lock
}

Thread C (priority 30) {
    // nothing
}
```

High priority threads should run to completion before those with lower priorities.
In this example, **B** gets blocked by the lower priority **A** which is holding the shared lock.
C has a higher priority than **A**, so once **C** starts running (without priority donation), **A** will not be able to run until **C** completes, at which point **A** will release the lock and **B** will have a chance to run.
Desired completion order is **B, C, A**.
Without priority donation, completion order is **C, B, A** ⊕

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

21

// src/tests/threads/priority-donate-single.c

```
static struct lock lock;

void test_priority_donate_single(void)
{
    thread_set_priority(20);
    lock_init(&lock);
    lock_acquire(&lock);

    thread_create("Thread B", 40, b_thread_func, NULL);
    thread_create("Thread C", 30, c_thread_func, NULL);

    lock_release(&lock);
    msg("Thread A finished.");
}

static void b_thread_func(void *aux UNUSED) {
    lock_acquire(&lock);
    lock_release(&lock);
    msg("Thread B finished.");
}

static void c_thread_func(void *aux UNUSED) {
    msg("Thread C finished.");
}
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

22

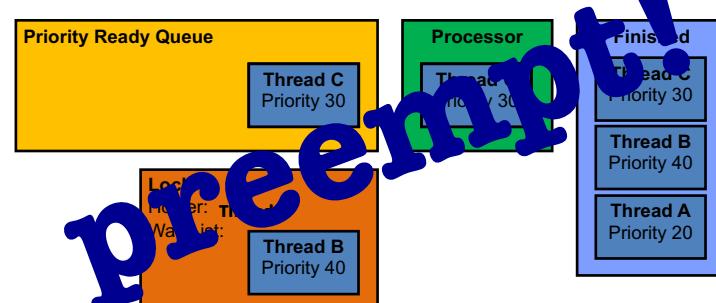
// src/tests/threads/priority-donate-single.ck

```
# -*- perl -*-
use strict;
use warnings;
use tests::tests;
check_expected ([<<'EOF']];
(priority-donate-single) begin
(priority-donate-single) Thread B finished.
(priority-donate-single) Thread C finished.
(priority-donate-single) Thread A finished.
(priority-donate-single) end
EOF
pass
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

23

Without Priority Donation



See class video and YouTube for animation

```
Thread A (priority 20) {
    acquire Lock
    create Thread B
    create Thread C
    release Lock
}

Thread B (priority 40) {
    acquire Lock
    release Lock
}

Thread C (priority 30) {
    // nothing
}
```

24

```
// src/tests/threads/priority-donate-single.ck
```

```
# -*- perl -*-
use strict;
use warnings;
use tests::tests;
check_expected ([<<'EOF']];
(priority-donate-single) begin
(priority-donate-single) Thread B finished.
(priority-donate-single) Thread C finished.
(priority-donate-single) Thread A finished.
(priority-donate-single) end
EOF
pass
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

25

Next Lecture

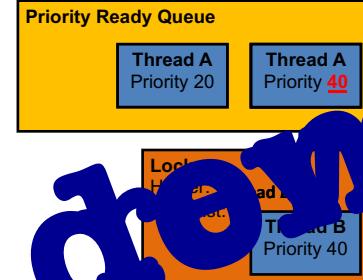
- Case Study: 4.4 BSD Unix
- Multi Core & Multi Processor Scheduling
- Lab 2 Secret Sauce



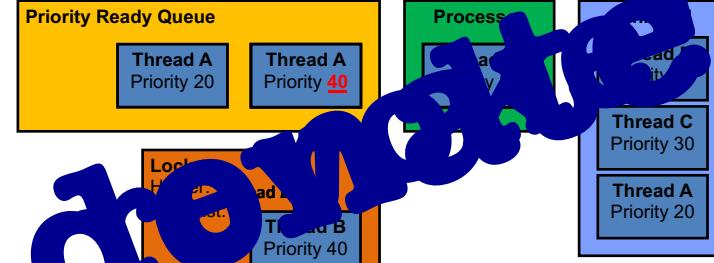
UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

27

With Priority Donation



See class video and YouTube for animation



```
Thread A (priority 20) {
    acquire lock
    create ThreadB
    create ThreadC
    release lock
}

Thread B (priority 40) {
    acquire lock
    release lock
}

Thread C (priority 30) {
    // nothing
}
```

26

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.