

PRINCIPLES OF COMPUTER SYSTEMS DESIGN

CSE130

Winter 2020

Scheduling III



Notices

- Lab 2 due 23:59 **Sunday February 9**

2

Today's Lecture

- Review of CPU Scheduling so far
- Case Study: 4.4 BSD Unix
- Multi Core and Multi Processor Scheduling
- Lab 2 Secret Sauce 

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

3

CPU Scheduling : Review

- Efficient use of the CPU
 - Fundamentally a pre-emptible resource
- Scheduling Opportunities
- Scheduler vs. Dispatcher
- Scheduling Criteria
 - Utilisation, Throughput, **Turnaround Time**, **Waiting Time**, Response Time
- Scheduling Algorithms
 - First Come First Served
 - Round Robin
 - Shortest Job First (Preemptive & Non-preemptive)
- Multi-level Queues
- Multi-level Feedback Queues
- Priority Donation

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

4

1

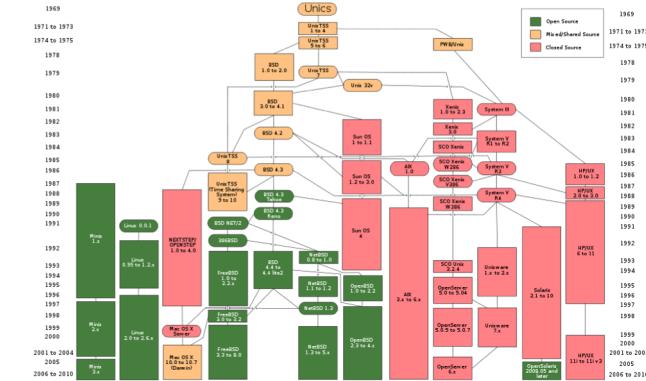
Pintos Scheduling : Review

- Pintos implements a single ready queue with a quantum
 - In Lab 2, you are implementing a priority scheduler ☺
 - There are 64 priorities, default is 31, 0 is lowest
 - Threads can change priorities
 - Every queue (ready, and any related to concurrency primitives like semaphores and condition variables) must reflect these priorities when removing (or adding) elements
 - **General principle is that no higher priority thread should wait while a lower priority thread executes**
 - **Preemption is super important, so get it right!**
 - But priority donation is problematic ☹

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

5

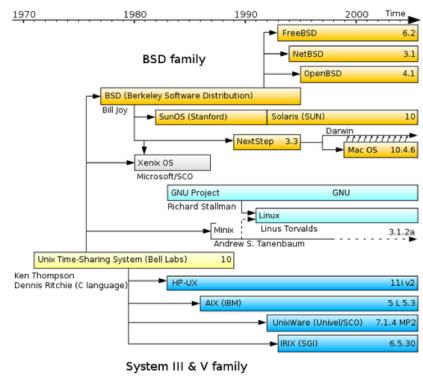
A Unix Family Tree



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

6

Another Unix family Tree



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David G. Harrison. All rights reserved.

7

Introduction to 4.4 BSD Unix

- **“Berkeley Software Distribution (BSD) is a Unix operating system derivative developed and distributed by the Computer Systems Research Group (CSRG) of the University of California, Berkeley, from 1977 to 1995”** - Wikipedia
 - BSD was given away free, but Berkeley was sued by AT&T for ripping off their intellectual property (AT&T Unix had to be paid for at the time and CSRG had a license) - AT&T won the case
 - First BSD released with no misappropriated AT&T code (4.4 BSD-Lite) was made available June 1994
 - **4.4 BSD-Lite Release 2** came out in 1995 and Berkeley shut down CSRG shortly thereafter
 - All current Apple Operating Systems are derived from 4.4 BSD

UCSC: BSOE CSE130 Winter 2020. Copyright © 2017-2020 David G. Harrison. All rights reserved.

8

4.4 BSD Unix CPU Scheduling

- Based on multilevel feedback queues
- Runnable threads explicitly assigned priority that determines which queue they're assigned to
- Always runs thread from highest priority, non-empty queue
- When a thread exhausts its quantum, it is returned to its original queue
- Quantum for all queues is 100ms
- User mode priority can be adjusted programmatically
- Threads move between queues when their priority changes

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

9

4.4 BSD Unix CPU Scheduling

- When a thread T_p other than the running thread T_c is promoted to a higher priority:
 - T_c is preempted and context switch to T_p is performed...
 - Unless T_c is in kernel mode (a system call has been made), then switch made to T_p as soon as T_c returns to user mode (i.e. the system call returns)
- Scheduling algorithm tailored to favour interactive processes
 - Priority of threads blocked on IO for 1 or more seconds (usually due to a slow human) have their priority increased ↑
 - Threads that use significant amounts of CPU have their priority decreased ↓

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

10

4.4 BSD Unix Priorities

- **Kernel mode priorities:**
 - Determined by activity
 - Smaller the number the higher the priority
 - In the range [0,49]
 - **User mode priorities:**
 - Can be set programmatically
 - But also manipulated by the kernel (see next slide)
 - In the range [50,127]
- | |
|------------------------------------------------|
| 0 - while swapping process (context switching) |
| 4 - waiting for memory |
| 8 - waiting for file ctrl information |
| 16 - waiting on disk i/o completion |
| 20 - waiting for kernel file system lock |
| 25 - PZERO, baseline |
| 24 - waiting on network socket |
| 32 - waiting for child to exit |
| 36 - waiting for user file system lock |
| 40 - waiting for an event/signal |
| 50 - PUSER, base user-mode priority |

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

11

4.4 BSD Unix : User Process Scheduling I

- Priorities in the range [50,127] (remember higher number is lower priority)
- Variables:
 - p_{usrpri} : user-mode scheduling priority of a process
 - p_{estcpu} : estimate of recent CPU utilization of a process
 - p_{nice} : user-settable weighing factor in range [-20,20]
 - p_{sptime} : time process spent sleeping in seconds
 - $load$: sampled mean of the sum of lengths of run queue and short-term sleep queue (e.g. processes waiting for disk I/O)
- User priority recalculated every 4 clock ticks (40ms):

$$p_{usrpri} \leftarrow 50 + p_{estcpu}/4 + 2 * p_{nice}$$
- Resulting value is capped into the range [50,127]

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

12

3

4.4 BSD Unix : User Process Scheduling II

- p_{estcpu} is incremented every time the clock ticks (10ms) and the process is executing
- In addition, once per second, CPU utilization of a runnable process is adjusted / recalculated:

$$p_{estcpu} \leftarrow (2 * \text{load} / (2 * \text{load} + 1)) * p_{estcpu} + p_{nice}$$

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

13

4.4 BSD Unix : User Process Scheduling III

- Example:
 - let T_i be the number of clock ticks accumulated during interval i
- Assume $\text{load} = 1$, starting from T_0 ...

$$\begin{aligned} p_{estcpu} &\leftarrow 0.66 * T_0 \\ p_{estcpu} &\leftarrow 0.66 * (T_1 + 0.66 * T_0) = 0.66 * T_1 + 0.44 * T_0 \\ p_{estcpu} &\leftarrow 0.66 * T_2 + 0.44 * T_1 + 0.30 * T_0 \\ p_{estcpu} &\leftarrow 0.66 * T_3 + \dots + 0.20 * T_0 \\ p_{estcpu} &\leftarrow 0.66 * T_4 + \dots + 0.13 * T_0 \end{aligned}$$

⇒ ~ 90% of previous CPU utilization forgotten after 5s
- For sleeping processes, p_{estcpu} is adjusted when the process wakes:

$$p_{estcpu} \leftarrow (2 * \text{load} / (2 * \text{load} + 1)) * p_{sliptime} * p_{estcpu}$$

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

14

Multi Core & Multi Processor Scheduling

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

15

Multi Core

- **Most application software does not need to do anything different on a multi-core CPU**
- Operating System assigns threads to cores
 - User code typically does not try to influence which core it runs on
- More cores may give **better performance** for:
 - Multi-threaded applications
 - Multi-tasking & time-sharing response times
 - Transaction processing
- However, the **speed up is non-linear** and in general exploiting parallelism in software can be tricky

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

16

Multi Processor / Multi Processing

- Scheduling gets more complex with multiple CPUs ☺
- Assume homogeneous processors (all Intel® Core™ i7, say) within a multiprocessor system, scheduling can be:
 - **Asymmetric (AMP)**
One CPU acts as a master and schedules all CPUs in the system
 - **Symmetric (SMP)**
Each core is self scheduling; if the ready queue is shared, care must be taken with synchronisation
- **SMP** supported by Windows, Linux, macOS, etc. etc.
- **AMP** is rare, sometimes used for loosely coupled systems
- Simultaneous Multi Threading (**SMT**) ([hyper-threading](#) on Intel®) is in some ways "on-chip" SMP

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

17

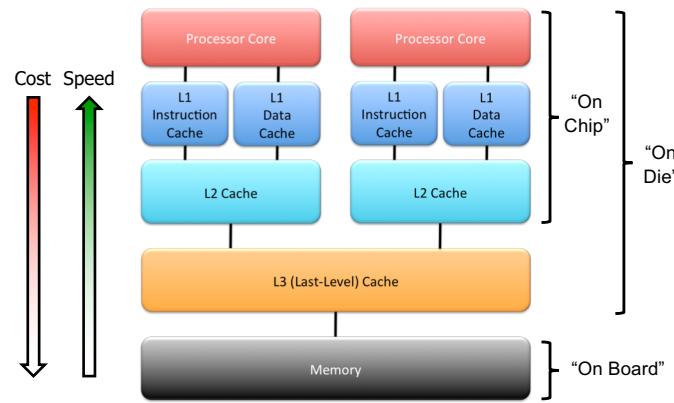
Load Balancing

- We can either have a single system-wide ready queue, or have one for each core / processor
- If the ready queue is not shared, SMP systems need to balance workload between the multiple queues:
 - **Push Migration:** If a core is overloaded, find a comparatively *under-loaded* core and migrate threads *to* its queue
 - **Pull Migration:** If a core is under loaded, seek an *over-loaded* core and migrate threads *from* its queue
- In practice, both schemes are used, depending on load
 - When lightly loaded, push
 - When heavily loaded, pull

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

18

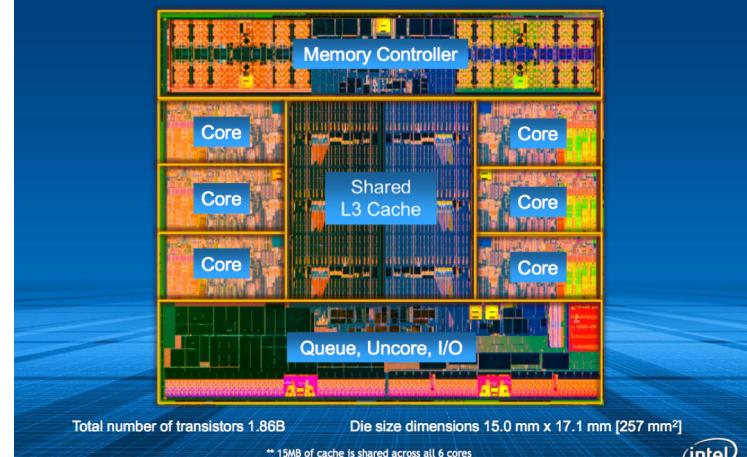
Modern CPU Memory Layout



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

19

Intel® Core™ i7-4960X Processor Die Detail



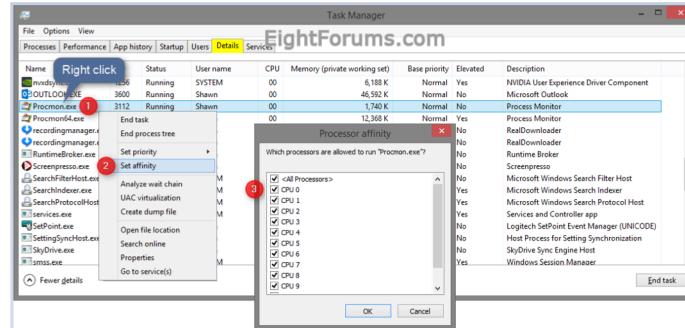
Hard Processor Affinity

- Scheduling a thread to a **different core than the one it previously ran** on will result in **L1** cache & **L2** cache misses, which results in increased **L3** cache hits
 - Also **L3** cache misses if multi-processor system
 - Potentially a significant performance hit ☹
- Hard Affinity** specifies where a thread is permitted to run
 - Implemented as a bit mask where each bit corresponds to a CPU
 - Defaults to all CPUs (i.e. no affinity)
 - Thread affinity mask must be subset of the containing process affinity mask, which in turn must be a subset of the active processor mask

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

21

Setting Hard Affinity - Windows® 10



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

22

Setting Hard Affinity - Linux

- Assume two processors
 - Each processor has four cores
 - All cores are hyper-threaded
 - i.e. `top` will show 16 CPUs ($2 \times 4 \times 2$)
 - or `cat /proc/cpuinfo | grep '^processor' | wc -l` returns 16
- Bind a process to the non-hyper threaded cores of the first processor:


```
$ taskset -c 0,2,4,6 ./myprogram
```
- Bind a process to all cores on the second processor:


```
$ taskset -c 8,9,10,11,12,13,14,15 ./myprogram
```

 or


```
$ numactl --cpunodebind=1 ./myprogram
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

23

Setting Hard Affinity - macOS

- No command line utilities provided ☹
- Affinity API exists, but is complex and not well documented
- OpenBSD (recall, macOS is also descended from BSD 4.4) has functionality equivalent to that found in Linux

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

24

Soft Processor Affinity

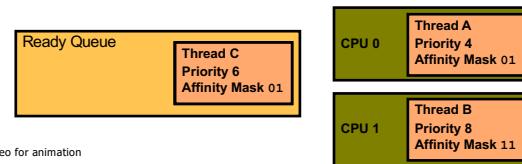
- **Soft Affinity** gives every thread an “ideal processor”
 - OS selects ideal processor for first thread in a process
 - Typically round robin across all CPUs
 - Next thread gets assigned a CPU relative to the process seed
 - Can be overridden by setting hard affinity
 - **OS attempts to return each thread to the CPU on which it was last resident** (i.e. the last one it executed some instructions on)

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

25

Affinity Collisions

- Highest-priority n threads (when n is No. of CPUs) may not be running if (soft or hard) thread affinity interferes
- Windows® guarantees highest-priority thread will be running ☺
 - But lower-priority $n-1$ ready threads may not be running ☹
 - Scheduler refuses to “move” running threads among CPUs if affinity has been set
- Example: (0 = highest priority)



See class video for animation

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

26

Lab 2 - Secret Sauce

- Must deal with Priority and Preemption
- Pintos’ Ready queue is a:
 - List of threads
- Semaphore waiting list is a:
 - List of threads
- Condition Variable waiting list is a:
 - List of semaphores
- Lock holder is a:
 - Thread



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

27

Priority Donation

```

Low Thread (priority 31) {
  acquire Lock A
  create Medium Thread
  create High Thread
  release Lock A
}
Medium Thread (priority 32) {
  acquire Lock B
  acquire Lock A
  release Lock A
  release Lock B
}
High Thread (priority 33) {
  acquire Lock B
  release Lock B
}

```

Low Thread priority = 31
 Medium Thread priority = 32
 High Thread priority = 33
 Lock A holder =
 Lock B holder =
 Current thread =
 High Thread Finished
 Medium Thread Finished
 Low Thread Finished

See lecture webcast for animation

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

41

Next Lecture

- Real-time Scheduling
- Scheduling Wrap-up

42

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.