

# PRINCIPLES OF COMPUTER SYSTEMS DESIGN

## CSE130

Winter 2020

CPU Scheduling IV



### Notices

- **Lab 2** due 23:59 **Sunday February 9**
- **Lab 1** Grades Released
  - Come see me if you got an unexpected grade
  - Take care copying lab installations
    - There's a hidden directory ( .pintos ) that needs be copied too
- **MIDTERM CANCELLED**
  - **Assignment 3** will now be written
  - Questions similar to final
  - Individual work, no conferring
  - Consider it an open-book midterm

2

### Today's Lecture

- Real-Time Scheduling
- CPU Scheduling Summary
- More Lab 2 Secret Sauce

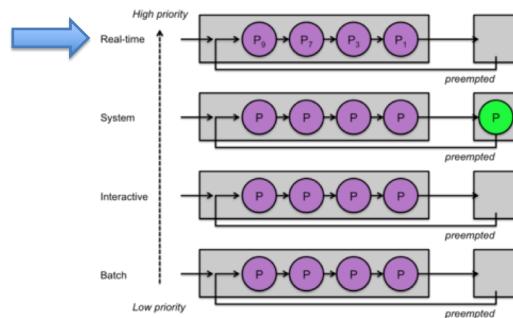


### Real-Time Scheduling

- **Hard Real-Time:** System guarantees to complete a critical task within a specified time period
- **Soft Real-Time:** Critical processes receive priority over less important ones, but no guarantees are given
- **Near Real-Time:** Marketing speak so manufacturers don't get sued if their Hard Real-Time OS were ever to not quite deliver 😊
  - IBM's Transaction Processing Facility (TPF) for their zSeries mainframes is a good example [https://en.wikipedia.org/wiki/Transaction\\_Processing\\_Facility](https://en.wikipedia.org/wiki/Transaction_Processing_Facility)
  - *"The depth of the CPU ready list is measured as any incoming transaction is received, and queued for the I-stream (processor) with the lowest demand, thus maintaining continuous load balancing among available processors."*

4

## Soft Real-Time Scheduling

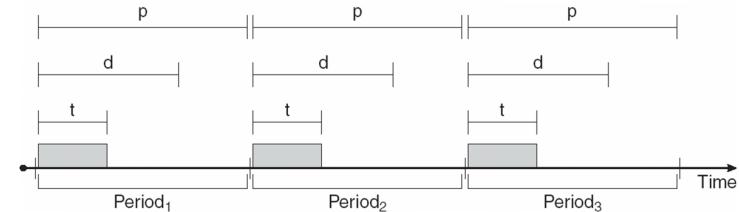


UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

Only one CPU - Diagram is not the best ☺ 5

## Real-Time Scheduling

- Periodic processes require the CPU for discrete periods
  - $p$  = duration of the **period**
  - $d$  = **deadline** by when the process must be serviced
  - $t$  = the processing **time**



- In most cases  $p = d$

6

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

## Can a Schedule be Found?

- If there are  $m$  periodic events and event  $i$  occurs with period  $P_i$  and requires  $C_i$  time units of CPU time to handle each event, then the load can be handled only if:

$$\sum_{i=1}^m \frac{C_i}{P_i} \leq 1$$

- **However:**

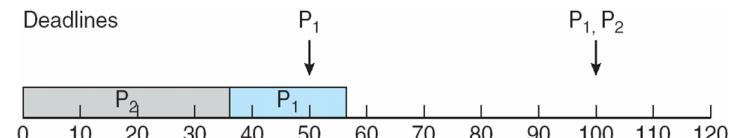
- Satisfying this equation does not mean a schedule exists, it means one might exist if you can find a suitable scheduling algorithm
- But failing to satisfy this equation absolutely means no schedule exists regardless of scheduling algorithm

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

7

## Deadline Unaware Scheduling

- We have 2 processes:
  - $P_1$  with  $p_1 = 50$ ,  $t_1 = 20$
  - $P_2$  with  $p_2 = 100$ ,  $t_2 = 35$



- If  $P_2$  runs before  $P_1$ , then  $P_1$  will miss its 1<sup>st</sup> deadline ☺
  - (assuming deadlines are at period ends)

8

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

## Rate Monotonic

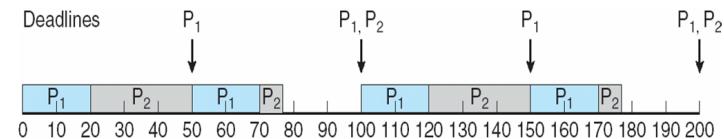
- **Monotonic:** A given order of process execution is preserved
- Processes assigned **priority inversely** based on their period
  - Shorter period = higher priority
  - Longer period = lower priority
- **Preemptive!**
  - If a long process does not finish within shorter processes period, longer process is kicked off the CPU
- Assumes:
  - CPU bursts are consistent
  - A deadline hitting schedule can be found

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

9

## Deadline Aware Scheduling

- With **Rate Monotonic Scheduling**,  $P_1$  has a higher priority because its period is shorter, again:
  - $P_1$  with  $p_1 = 50, t_1 = 20$
  - $P_2$  with  $p_2 = 100, t_2 = 35$
- Now  $P_1$  will preempt  $P_2$  at  $T = 50$



- And deadlines for both  $P_1$  and  $P_2$  are met ☺

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

10

## Rate Monotonic

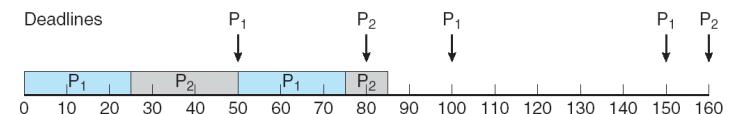
- **Considered optimal:**
  - If a set of processes cannot be scheduled by Rate Monotonic, the set cannot be scheduled by any other algorithm that assigns ‘static’ priorities
- As more processes need to be scheduled, attainable system utilisation falls

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

11

## Missed Deadlines with Rate Monotonic

- If  $P_1$ ’s processing time is increased,  $t_1 = 25$  (was 20)
- And  $P_2$ ’s period is decreased,  $p_1 = 80$  (was 100)
- $P_1$  still has a higher priority because its period is less than  $P_2$



- $P_1$  will preempt  $P_2$  at  $T = 50$  and  $P_2$  will miss first deadline ☹
  - If  $P_2$  had not been preempted by  $P_1$ , both deadlines would be met
  - i.e. RM unable to find a schedule, even though CPU demand is <100%

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

12

## Earliest Deadline First

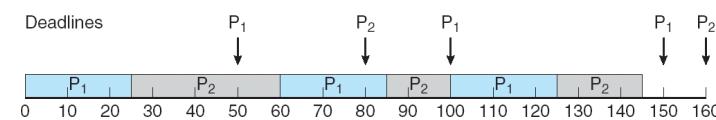
- Priorities are assigned dynamically according to deadlines
  - Earlier deadline => higher priority
  - Later deadline => lower priority
  - Equivalent to preemptive SJF**
- When put on ready queue, process states its deadline
  - Unlike SJF, we know the deadline, no guesses needed ☺
- The process with the earliest deadline is run first
- Dynamic, so can cope with:
  - Non periodic processes**
  - Variable processing times**
    - Processes can change their mind about their deadline each time they enter ready state (i.e. each time they get onto the Real-Time Ready Queue )

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

13

## Earliest Deadline First: Example

- Same schedule that Rate Monotonic failed:
  - P1**  $p_1 = 50$
  - P1** runs and finishes, **P1**'s first deadline met
  - P2** runs, **P1** does not preempt as **P2**'s first deadline earlier than **P1**'s second deadline
  - P2** runs and finishes, **P2**'s first deadline met
  - and so on...



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

14

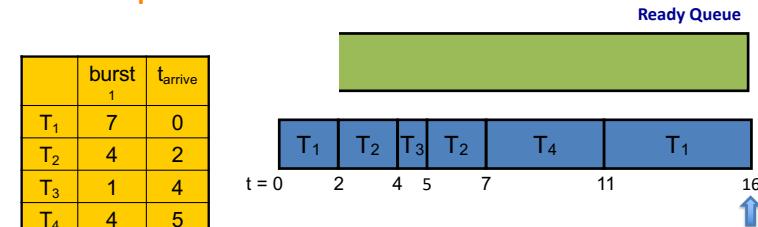
## CPU Scheduling Summary

- Efficient use of the CPU** ( fundamentally a preemptible resource )
- Scheduling Opportunities & Scheduler vs. Dispatcher**
- Scheduling Criteria** ( Utilisation, Throughput, Turnaround Time, Waiting Time, Response Time )
- Scheduling Algorithms** ( FCFS, RR, SJF, Priority Based, Preemptive & Non-preemptive)
- Multi-level Queues & Multi-level Feedback Queues**
- Real World Example : **4.4 BSD Unix**
- Priority Inversion** ( the problem ) **Priority Donation** ( the solution )
- Multi Core and Multi Processor** ( L1/L2/L3 Caches, AMP, SMP, Push/Pull, Processor Affinity)
- Real-Time Scheduling** ( Types, Schedulability, RM, EDF )

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

15

## Preemptive SJF



$$T_1 \text{ waiting time} = (0-0) + (11-2) = 9$$

$$T_2 \text{ waiting time} = (2-2) + (5-4) = 1$$

$$T_3 \text{ waiting time} = 4-4 = 0$$

$$T_4 \text{ waiting time} = 7-5 = 2$$

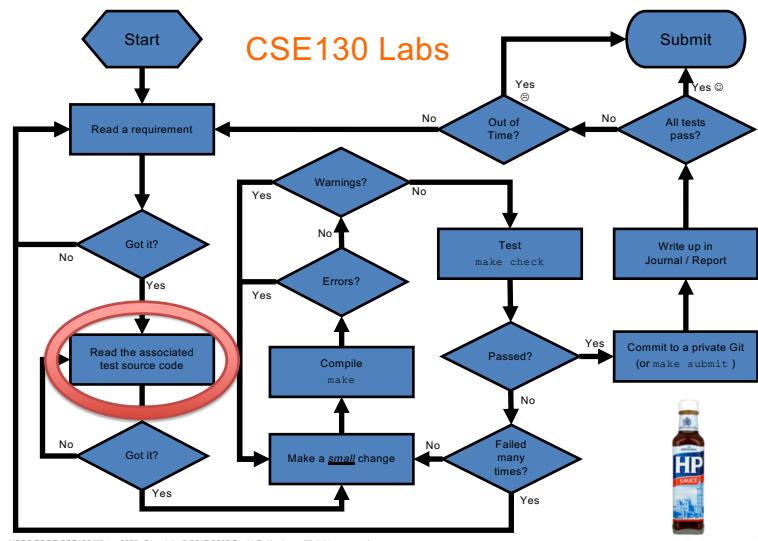
See Class Webcast for Animation

$$\text{Mean Waiting Time} = (9 + 1 + 0 + 2) / 4 = 3$$

$$\text{Mean Turnaround Time} = ((16-0) + (7-2) + (5-4) + (11-5)) / 4 = 7$$

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

32



33



## Lab 2 - Secret Sauce

- Read the tests line-by-line
  - How many threads are there?
  - What state(s) are they in?
  - Which one is on the CPU?
  - What lists are the others in?
- Draw Pictures
  - Use the whiteboards
- Discuss with classmates
  - But do NOT share code

34

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

## Lab 2 - Secret Sauce

- Must deal with Priority and Preemption
- Pintos' Ready queue is a:
  - List of threads
- Semaphore waiting list is a:
  - List of threads
- Condition Variable waiting list is a:
  - List of semaphores
- Lock holder is a:
  - Thread
- Once you get into the “extreme” requirements
  - The TAs and I will offer only minimal guidance
  - Priority will be given to students stuck on earlier tests



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

35



## Lab 2 - Secret Sauce

- General Approach
  - Only write code to pass the test you’re working on right now
  - Once you’ve passed the test, make submit to “bank” that grade
- Common mistakes
  - priority-sem
    - Yielding before incrementing the semaphore value
  - priority-condvar
    - Trying to sort the waiting list as if it were a list of threads
  - priority-donate-single
    - How many locks can one thread hold?
    - Is the lock being released the one that triggered the donation?
  - All tests
    - Dealing with priority but not preemption

36

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.

## Next Lecture

- Introduction to Operating System Security

37

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All rights reserved.