

PRINCIPLES OF COMPUTER SYSTEMS DESIGN

CSE130

Winter 2020

Memory Management II - Allocation



Notices

- **Lab 3** due **Sunday March 1**
- **Assignment 3** due **Sunday February 16**
 - Equivalent to two questions in the final
- My office hours changed today only
 - 3pm to 4pm
- No class on Monday February 17
 - President's Day

2

Today's Lecture

- Brief Recap
- Memory Allocation
 - Partitioning Schemes
 - Fragmentation
 - Allocation Schemes
 - Paging
 - Paging on the Intel 80386
- Lab 3 - Secret Sauce



Operating Systems & Memory

- Computers consist of **CPUs**, **memory**, and **I/O devices**
- An Operating System's primary concerns are **equitable CPU scheduling**, and **safe and secure memory management**
- This series of lectures largely ignores the on-chip L1 & L2, and on-die L3 caches to concentrate on the off-chip, but on-board, Random Access Memory (RAM)
 - However, L cache misses still have to be taken into consideration in CPU scheduling algorithms, so don't forget about the L caches
- Memory Types:
 - Registers fastest tiny bytes
 - L Caches fast small KB to MB
 - **RAM** **slow** **large** **GB**
 - Hard Disk very slow massive TB

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

3

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.


4

1

Memory Protection

The relocation-register scheme:

- **Relocation Register**
 - contains the smallest permitted **physical** address
- **Limit Register**
 - contains the maximum **logical** address



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

5


Memory Partitioning

- Main memory must hold the code and data for both the OS and all user processes
- Main memory needs to be **divided into partitions**:
 - The resident operating system, usually held in low memory
 - User processes held in high memory
- Typically we want several user processes simultaneously in memory (multiprogramming / time-sharing)
- Need to partition memory to protect processes from each other
- Also must **allocate memory** to processes in a **safe & timely** manner

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

6

Memory Allocation



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.


Questions:

- When is memory allocated?
- How much memory?
- Which memory?
- Contiguous memory?
- Can process be swapped?

7

Fixed Partitions

- Traditional first, simple approach
- Memory statically partitioned
- Partitions of different, fixed sizes
- Processes allocated **smallest available partition large enough for their needs**
- Addresses bound at load time
- Fixes partition in which process runs
- **Problem?**
 - **Internal fragmentation** (wasted space)




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

8

2

Variable Partitions



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

9

Fragmentation

- **External Fragmentation**
 - Total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation**
 - Allocated memory may be slightly larger than requested memory
 - This size difference is memory internal to a partition, but not being used

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

10

Allocation Algorithms

How to satisfy a request from a list of free holes?

- **First-fit**
 - Allocate the first hole large enough
- **Best-fit**
 - Allocate the smallest hole large enough
 - Must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- **Worst-fit**
 - Allocate the largest hole available
 - Must also search entire list
 - Produces the largest leftover hole

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

11

External Fragmentation

We can reduce external fragmentation via:

- **Compaction**
 - Shuffle memory contents around to place all free memory together in one large block
 - Compaction is possible only if relocation is dynamic, and is done at execution time
 - **Problem?**
 - All user processes have to wait whilst OS re-shuffles the free memory
- **OR...**
 - Map logical address space of a process to non-contiguous physical memory
 - Process is allocated physical memory from where ever it is available
 - This is known as **PAGING**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

12

Paging: Definitions

- **Page** (verb) “Divide a piece of software or data into sections, keeping the most frequently accessed in main memory and storing the rest in virtual memory” – Google
- **Paging** (noun) “The transfer of pages of data between the main memory of a computer and its auxiliary memory” – Collins English Dictionary
- **Paging** (noun) “A technique of storage management that transfers pages from secondary storage to main storage when they are required, and returns them to secondary storage when they are not” – dictionary.com

Paging for Memory Management

- **Paging** (the concept) is usually what people are talking about today when they say “**virtual memory**”
- Virtual memory requires hardware and software support
 - Hardware component is the **MMU** (memory management unit)
 - Operating System decides **how to translate virtual <=> physical addresses**


UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

13

14

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.


Paging Overview



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

15

Paging Overview



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

16

Paging Overview

- Based on the notion of a virtual address space
 - The **illusion** of a large, contiguous address space
 - Virtual address space >> Physical address space**
 - Each **process** gets its own virtual address space
 - Threads do not get their own
- Divides the address space into fixed size **pages**
 - Virtual Page:** A “chunk” of the virtual address space
 - Physical Page:** A “chunk” of the physical address space
 - Also known as a frame
 - Size of virtual page == size of physical page**

Paging Overview

- We need to **MAP** virtual pages to physical pages
 - By itself, a virtual page is an illusion
 - Cannot actually store anything
 - Needs to be backed by a physical page
 - Before a virtual page can be accessed...
 - It must be paired with a physical page
 - i.e. it must be **mapped** to physical page
 - This mapping must be stored somewhere
 - On every access to the virtual page...
 - Its mapping is looked up
 - Access is then directed to the relevant physical page


UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

17

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.


18

Paging Overview



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

19



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

20


Paging in the Intel 80386

- Intel 80386 (Mid 1980's)
 - 32-bit processor
 - 4KB page size
- What is the **virtual address space**? (the addressable memory)
 - $2^{32} = 4\text{GB}$
- How many **virtual pages** in the virtual address space?
 - $4\text{GB}/4\text{KB} = 1\text{M}$
- What is the size of the **physical address space**?
 - It depends, but $\leq 4\text{GB}$ (i.e. no point in installing $> 4\text{GB}$ memory as it can't be addressed)
- How many **physical pages** are there?
 - It depends, but $\leq 1\text{M}$
- Even though physical address space may be less than 4GB, we'll assume physical addresses are 32bits

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

21

Intel 80386: Virtual Pages



Highest 20 bits are the virtual page number
Lowest 12 bits are the offset within that page


See class video for animation

22

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

Intel 80386: Address Translation

- Assume Virtual Page 7 is mapped to Physical Page 34
- We want to access some offset in Virtual Page 7





UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

23

Intel 80386: Address Translation

- Assume Virtual Page 7 is mapped to Physical Page 34
- We want to access some offset in Virtual Page 7




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

24

Intel 80386: Address Translation

- Assume Virtual Page 7 is mapped to Physical Page 34
- We want to access some offset in Virtual Page 7



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

25

26

A Note on Notation

- We **MAP** pages
 - Specifically, virtual pages to physical pages
- We **TRANSLATE** addresses
 - Specifically, virtual addresses to physical addresses

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

Intel 80386: VPN->PPN

- Problem:** Need to keep track of VPN -> PPN mappings
 - VPN 65 -> PPN 981
 - VPN 3161 -> PPN 1629
 - VPN 9327 -> PPN 524

- Solution: Page Table**

- A “lookup table” for the mappings
- Can be thought of as an array
- Each array element is called a **page table entry** (PTE)

```
uint32_t PAGE_TABLE[1<<20];
PAGE_TABLE[65] = 981;
PAGE_TABLE[3161] = 1629;
PAGE_TABLE[9327] = 524;
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

27

28

Intel 80386: Page Tables

Two problems with page tables:

- Page table is too large**
 - Page table has 1M entries
 - Each entry is 4B (20-bit PPN)
 - Page table is 4MB** (ouch!)
 - Very expensive in the 1980s
 - Still problematic today on embedded devices

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

Intel 80386: Page Table Too Large

- Typically, the majority of PTEs are empty

```
PAGE_TABLE[0] = 141;
...
PAGE_TABLE[532] = 1190;
```

Unallocated (i.e. wasted) space!

```
PAGE_TABLE[1048402] = 845
...
PAGE_TABLE[1048575] = 742; // 1048575 == (2^20) - 1
```

- Unallocating every empty PTE is tedious
 - Instead, only unallocate long sections of empty PTEs

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

29

Intel 80386: Page Table Too Large

- How do we allow PTEs to become “unallocated”?
 - The page table must be restructured

- Before restructuring: **flat**

```
uint32_t PAGE_TABLE[1024*1024];
```

- After restructuring: **hierarchical**

```
uint32_t *PAGE_DIRECTORY[1024];
PAGE_DIRECTORY[0] = NULL; // 1024 PTEs unallocated
PAGE_DIRECTORY[1] = NULL; // 1024 PTEs unallocated
PAGE_DIRECTORY[2] = malloc(sizeof(uint32_t)*1024);
PAGE_DIRECTORY[2][0] = 753;
PAGE_DIRECTORY[2][1023] = 21;
PAGE_DIRECTORY[3] = NULL; // 1024 PTEs unallocated
PAGE_DIRECTORY[4] = NULL; // 1024 PTEs unallocated
```

30

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

Intel 80386: Page Tables

Two problems with page tables:

- Page table is too large
 - Page table has 1M entries
 - Each entry is 4B (20-bit PPN)
 - Page table is 4MB (ouch!)
 - Very expensive in the 80s
 - Still problematic today on embedded devices
- Page table is stored in memory**
 - Before every memory access, we always have to fetch the PTE from comparatively slow RAM => big performance penalty

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

31

Intel 80386: Accelerating Translation

- Problem:** Retrieving PTEs from memory is slow ☹

- Solution:** “Cache” PTEs inside the processor ☺

- Translation Lookaside Buffer (TLB)

- “Lookaside Buffer” is an archaic term for a cache

- Whenever a virtual address needs to be translated, the TLB is searched:
 - hit vs. miss (we'll look at these next week)

- 32-entry TLB on 80386

- Each TLB entry has a **tag** and some **data**

- Tag: 20-bit VPN + 4-bit metadata

- Data: 20-bit PPN

32

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

Next Lecture

- Address Translation
- TLB Hits and Misses
- Demand Paging

39

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.