

PRINCIPLES OF COMPUTER SYSTEMS DESIGN

CSE130

Winter 2020

Memory Management V - Frame Allocation



Notices

- Lab 3 due **Sunday March 1**
- Assignment 4 due **TODAY**

2

Today's Lecture

- Page vs. Frame
- Bélády's Anomaly
- Page Replacement Algorithms
- Frame Allocation Algorithms

- Lab 3 - System Calls Secret Sauce 

Naming Confusions 😞

- **Page Fault**
 - In the TLB:
 - A Page Table entry was not found
 - In the Page Table:
 - A Virtual Page does not currently map to a Physical Page
- **Page v.s. Frame**
 - Virtual Page = Page
 - Physical Page = Frame (really confusingly, also known as a "Page Frame")
- **Important Algorithms**
 - Page Replacement
 - Frame Allocation

Memory Management

- Process memory requirements tend to grow and grow
 - Each page fault requires the allocation of a frame
 - Eventually we run out of frames ☺
- **When / How do we reclaim memory?**
 - **Process terminates**
 - All frames reclaimed
 - **Swap entire processes out**
 - Some or all frames reclaimed
 - **Swap out some pages from an active process**
 - Some frames reclaimed

Page Replacement

- Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement
- Use **modified** (dirty) **bit** in page table entry to reduce overhead of page transfers - only modified pages are written to disk
- Page replacement completes the separation between logical memory and physical memory - **a large virtual memory space can be provided on a much smaller physical memory**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

5


UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

6

Page Replacement & Frame Allocation

- **Page-Replacement** algorithm
 - Wants lowest page-fault rate both on first access and subsequent re-access
 - Decides which pages to replace
- **Frame-Allocation** algorithm
 - How many frames should we give each process?
 - Which frames should be replaced?
- We evaluate algorithms by running them on a **particular sequence of memory references** and computing the number of page faults
 - Sequence is just page numbers, not full addresses
 - Repeated access to the same page does not cause a page fault

Page Faults vs. Available Frames



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

7

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

8

2

Page Replacement Algorithm : FIFO

- Simple **First-In-First-Out** Queue
 - Once queue is full, page at the front of the queue is evicted on page fault
 - New pages are added at the back of the queue

Page Request	7	0	1	2	0	3	0	4
Queue	7	7 0	7 0 1	2 0 1	2 0 1	2 3 0	2 3 0	4
Page Faults	1	2	3	4	4	5	6	7

Bélády's Anomaly

Page Request	3	2	1	0	3	2	4	3	2	1	0	4
Page Faults	3	3 2	3 2 1	0 1	3 1	0 3 2	4 2	4 3	4 2	4 1	0 0	4 0
Page Faults	1	2	3	4	5	6	7	7	7	8	9	9!

Adding more frames does not necessarily result in fewer page faults ☺

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

9

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

10

Page Replacement Algorithm : OPT

- Optimal Page Replacement**
 - Replace page that will next be used **furthest in the future**
 - Currently impossible to implement ☺
 - But can give us a theoretical benchmark against which to assess alternatives



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

11

Page Replacement Algorithms : LRU & LFU

- Least Recently Used**
 - Use past rather than future knowledge
 - Replace page that has not been used for the longest time
 - Associate **time of last use** with each page
- Least Frequently Used**
 - Like LRU, use past knowledge
 - Replace page that has been used least **over a given time period**
 - Associate **usage count** with each page

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

12

Stack Based Page Replacement Algorithms

- LRU, LFU, and OPT are stack algorithms **without** Bélády's Anomaly
- **Stack Algorithms**
 - Keep a stack of page numbers in a double linked list
 - When page referenced:
 - Move it to the top of the stack (head of the list)
 - 6 pointers need to be changed (inexpensive)
 - But each update more expensive
 - No search for replacement

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

13

LRU Approximations


- **Problem:**
 - LRU tends to need hardware assistance and can still be slow
- **Solution:**
 - Add a reference bit to each page
 - Initial value = 0
 - When page is referenced bit set to 1
 - Replace any pages where reference bit = 0 (we're obviously not using them)
- **Problems?**
 - We **do not know the order** in which the referenced pages were accessed
 - What happens when all pages have been referenced?

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

14

LRU Approximations

- **Second Chance**
 - Extends FIFO
 - Each page gets two chances
 - On second access, page goes to back of queue
 - Frequently referenced pages don't get replaced quite so often
- **Circular Queue**
 - Pointer indicates next page to replace
 - Advance until page with 0 reference bit found
 - Clear reference bits as you go
 - Replace victim with new page and set reference bit to 1
 - Reference bit also set to 1 when accessed



See class video for animation

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

15

Counting Page Replacement Algorithms

- Every page table entry has a **counter**; every time page is referenced through this entry, the counter is incremented
- When a page needs to be replaced, examine the counters to find **smallest value**
- But... linear search through table required ☺
- **Least Frequently Updated**
 - Replaces page with the smallest count
- **Most Frequently Updated**
 - Replaces page with the largest count

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

16

Clock Based Page Replacement Algorithms

- Every page entry has a **timestamp**
- Every time page is referenced, copy the clock into the timestamp
- When a page needs to be replaced, look at the timestamps to find **smallest value**
- But... linear search through table required ☺

Page Buffering Algorithms

- Idea: **Keep a pool of free frames**
 - A frame is “always” available when needed, not searched for at fault time
 - Read page into free frame and select victim to evict
 - When convenient, evict victim and add recovered frame to free pool
- Enhancement: **Keep list of modified pages**
 - When backing store otherwise idle, write pages there and set to non-dirty
- Enhancement: **Keep free frame contents intact**
 - If referenced again before frame is reused, no need to re-load contents from disk
 - Generally useful to reduce penalty if wrong victim frame selected

17

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

18

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

Frame Allocation

- Each process needs a **minimum number of frames**
 - Consider our pseudo machine instruction `ADD C, A, B`
 - Adds the integer at address A to the integer at address B, saves the result at address C
 - Lets say the instruction itself is 6 bytes and could span 2 pages
 - Possibly need 2 pages to handle each of A, B, and C
 - Possible 8 pages required, so allocate 8 frames
- Maximum is total number of frames in the system
- Two major allocation schemes
 - **Fixed allocation**
 - **Dynamic allocation**
- Many variations

Frame Allocation

- **Equal allocation**
 - E.g. If there are 93 frames (after allocating frames for the OS) and 5 processes, give each process 16 frames
 - i.e. keep some as a free-frame pool
- **Proportional Allocation (by size)**
 - Allocate according to the size of each process
 - Dynamic as the degree of multiprogramming and process sizes change
- **Proportional Allocation (by priority)**
 - If process P generates a page fault,
 - Instead of selecting one of P's pages for replacement
 - Select a page from process with lower priority and allocate the victim's freshly de-allocated frame to P (P ☺ but lower priority process ☺)
 - **Is this a reasonable idea?**
 - **On SMP and/or multi-core systems, how might you modify the algorithm?**

19

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

20

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

Global vs. Local Allocation

- **Global Allocation**
 - Process selects a replacement frame from the set of all frames
 - One process can take a frame from any another
 - Remember, memory is fundamentally preemptable
 - But process execution time can vary greatly
- **Local Allocation**
 - Each process selects from only its own set of allocated frames
 - More consistent per-process performance
 - But possibly underutilized memory
- Global allocation leads to better throughput than local allocation at the expense of unpredictable performance for individual processes

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

21


Thrashing Revisited

- When a process accesses a physical page that was swapped out, the page is swapped back in to physical memory
 - May cause another physical page to be swapped out
 - If this happens too rapidly, it's known as **THRASHING**
 - Extreme performance degradation ☹
- If a process does not have "enough" pages, the page-fault rate can be very high indeed
 - Page fault raised to get desired page
 - Replaces existing page
 - But quickly need replaced page back again
 - Leads to Low CPU utilization...
 - Operating system thinks that it needs to increase the degree of multiprogramming
 - Another process gets added to the system with negative results ☹

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

22

Thrashing



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

23

Preventing Thrashing


- **Problem:**
 - Processes stealing frames from one another leads to thrashing ☹
- **Solution:**
 - Only allow processes to trash themselves
 - Still affects other processes through contention for disk ☹
- **Better Solution?**
 - Observe that most processes access clustered sets of pages at a time
 - Use this to work out optimal per-process allocation

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

24

Locality Model

- A set of frames in frequent use by a process is known as a **locality**
- A process usually consists of a number of distinct localities
- As a process executes they tend to move from locality to locality
- Localities may overlap




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

25

Working Set Model

- Dynamically adjusts frames allocated during the previous τ
- $W_t(\tau)$ = set of pages referenced in $(t-\tau, t)$
- $w_t(\tau) = |W_t(\tau)|$ = working set size at time t



- Problems?

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

26

Keeping Track of the Working Set


- Approximate with **interval timer Δ + reference & history bits**
 - Timer interrupts every $\Delta/2$ time units
 - Keep in memory **2 history bits** for each page
 - On timer interrupt - **set all reference bits to 0**
 - Reference bit set or one of the history bits set => page is in working set
- Example: $\Delta = 10,000$
 - @ 0 : reset reference bits to 0 and history bits to 00
 - @ 5,000 : if reference bit = 1, set first history bit (i.e. bits are 01)
 - @ 10,000 : if reference bit = 1, set second history bit (i.e. bits now 10 or 11)
 - On page fault:
 - If history bits are 01, 10, or 11, or reference bit set, page is in working set
- Problems?
 - It's a coarse grained approach, a tradeoff between performance and accuracy

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

27

Track Page Fault Frequency

- More direct approach than the Working Set Model
- Establish "acceptable" page-fault frequency rate and use a local replacement policy



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

28

Page Replacement & Frame Allocation

- **Page v.s. Frame**
 - Virtual Page = Page
 - Physical Page = Frame
- **Page-Replacement** algorithm
 - Wants lowest page-fault rate both on first access and subsequent re-access
 - Decides which pages to replace
- **Frame-Allocation** algorithm
 - Decides how many frames to give each process
 - Decides which frames to re-allocate when necessary

Next Lecture

- File Systems I - Basics

29

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

34

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.