

# PRINCIPLES OF COMPUTER SYSTEMS DESIGN

## CSE130

Winter 2020

### File Systems II - The Virtual File System



### Notices

- Lab 3 due **Sunday March 1**
- Assignment 5 will be available **Monday March 2**
- Week 10 Lectures ( Final Prep ) will not be webcast

2

### Today's Lecture


- File System Basics
- The Layered File System
- Logical to Physical Mapping
- The Virtual File System
- More Lab 3 Secret Sauce 

### File Systems : The Basics

- Hard Disk Drives ( HDD )
- Solid State Drives ( SSD )
- Logical to Physical Mapping
- The Virtual File System
- Partitions
- Mount Points

## How Does a Hard Drive Work?

- Key Attributes
  - Read / Write any sector
  - Direct access
  - Supports sequential, indexed etc.
  - **Seeks are slow** (mechanical)
- To read/write a disk sector:
  - Move head to cylinder (seek)
  - Wait for sector to rotate to head (latency)
  - Read or write (transfer)


UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

5

## How Does a Solid State Drive Work?

- **RAM is Volatile Memory** - Information is lost when power removed
- **Flash Memory is Non Volatile**
  - NOR: Transistors wired in parallel
  - NAND: Transistors wired in series
    - Less expensive than NOR
    - Two transistors to a "cell"
    - Single Level Cell (SLC) can store one bit per cell ( 0 or 1 )
    - **Multi Level Cell (MLC)** store two bits per cell ( 00, 01, 10, or 11 )
      - Inexpensive, better storage capacity, but "wears out" faster than SLC
- **MLC NAND Flash Memory**
  - Typically can deliver 50,000 write/erase operations per cell during it's life
  - Device drivers & disk controllers implement "wear leveling" to make sure individual cells don't "wear out" faster than any other

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

6

## HDD vs. SSD (a non scientific evaluation)

- **Program Load**
  - SSD twice as fast as HDD
- **Data Read/Write**
  - SSD ten times as fast as HDD
- **Power Consumption**
  - SSD consume less power than HDD
  - SSD generates less heat
  - SSD produces less noise
- **Price**
  - SSD around five times more expensive than HDD for the same capacity in the 128GB to 500GB range
- **Lifespan**
  - HDDs are expected to last decades, SSDs just years

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

7


## HDD & SSD Working Together (non scientific)

- **SSDs Store:**
  - Things that get read a lot, are small, and you don't have many of
    - The Operating System
    - User Programs
- **HDDs Store:**
  - Things that get written a lot
    - Transaction Databases
    - User Documents (Word, Excel, Source Code)
    - Memory Swap Space
  - Things that are large
    - Movies
    - Genome Sequencer Output
  - Things you have lots of
    - Music
    - Images

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

8

## The Layered File System




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

9

## Logical to Physical Mapping

- File is a sequence of **Logical Records**
- Basic file system is an array [0..N] of **disk blocks**
  - Multiple disks => multiple arrays
- Mapping function required to go from logical to basic file system




- We need to store which block on which disk each file resides in
- **Problems?**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

10

## The Virtual File System

- Goal: **Seamlessly integrate different file systems**
  - File IDs unique within partition
  - Hide different semantics for different physical file system types
  - Provide network transparency
    - User program does not know if the file system it is using is local or remote
- **Virtual File System**
  - Supports generic interface
  - Identifiers unique network wide



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

11

## File System Structures

- **Boot control block**
  - Dedicated to OS for booting
  - The first block of a volume
- **Partition control block**
  - Partition - an instance of a file system
    - Multiple partitions per disk, or
    - One partition spanning multiple disks
  - PCB contains details of this instance
- **Directory structure**
- **File control blocks**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

12

## Disks & Partitions

```
$ diskutil list
/dev/disk0 (internal, physical):
#           TYPE NAME      SIZE    IDENTIFIER
0:  GUID_partition_scheme          *1.0 TB   disk0
1:      EFI EFI        209.7 MB  disk0s1
2:      Apple_HFS Macintosh HD    999.3 GB  disk0s2
3:      Apple_Boot Recovery HD    650.0 MB  disk0s3

/dev/disk1 (external, physical):
#           TYPE NAME      SIZE    IDENTIFIER
0:  GUID_partition_scheme          *2.0 TB   disk1
1:      EFI EFI        209.7 MB  disk1s1
2:      Apple_HFS Time Machine    2.0  TB   disk1s2
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

13

## Partitions, Mounting and Mount Points

- **Partition** contains raw or cooked (formatted) file system
- Boot information has own partition (boot partition)
- **Root partition** contains the OS
- Other partitions can hold other OSes, other file systems, or be raw
- A file system must be **mounted** before it can be accessed
- A file system is mounted at a **mount point**
  - Windows automatically mounts all volumes
  - Unix mounts some but not others
- Mounting creates entry in a **mount table**
- At mount time, file system **consistency** checked
  - If inconsistent, (try to) fix it, attempt mount again
  - If yes, add to mount table, allow access

15

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

## Mount Points : df

```
$ df -P -m
Filesystem 1M-blocks  Used Available Capacity Mounted on
/dev/disk0s2  953049  294823   657976  31% /
devpts       0       0       0  100% /dev
map -hosts   0       0       0  100% /net
map auto_home 0       0       0  100% /home
/dev/disk1s2 1907401 1240123  667278  66% /Volumes/Time Machine
```

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

16

Filesystem	1M-blocks	Used	Available	Use%	Mounted on
udev	1954	0	1954	0%	/dev
tmpfs	395	41	354	11%	/run
/dev/sdal	18915	16905	173	99%	
tmpfs	1974	0	1974	0%	/dev/shm
tmpfs	5	1	5	1%	/run/lock
tmpfs	1974	0	1974	0%	/sys/fs/cgroup
cgmfs	1	0	1	0%	/run/cgmanager/fs
bsd-05.soe.ucsc.edu:/export/home/sentinel	5242880	2276533	2966347	44%	/soe/sentinel
bsd-05/export/local/share	51200	3800	30395	8%	/share/local
bsd-05/export/local/share/software	101200	74812	38375	72%	/share/software
bsd-05.soe.ucsc.edu:/export/cse/classes	512400	3874	48177	8%	/share/share
bsd-05.soe.ucsc.edu:/export/home/dcharris	5242880	2276533	2966347	44%	/cse/classes
bsd-05.soe.ucsc.edu:/export/home/heidi2	5242880	2276533	2966347	44%	/cse/heidi2
bsd-05.soe.ucsc.edu:/export/home/gao15	5242880	2276533	2966347	44%	/cse/gao15
bsd-05.soe.ucsc.edu:/export/home/jmustain	5242880	2276533	2966347	44%	/cse/jmustain
bsd-05.soe.ucsc.edu:/export/home/pohn	5242880	2276533	2966347	44%	/cse/pohn
bsd-05.soe.ucsc.edu:/export/home/public	51240	1821	9420	18%	/soe/public
bsd-05.soe.ucsc.edu:/export/home/xinli	5242880	2276533	2966347	44%	/cse/public
bsd-05.soe.ucsc.edu:/export/home/awl	5242880	2276533	2966347	44%	/soe/awl
bsd-05.soe.ucsc.edu:/export/home/tjg	5242880	2276533	2966347	44%	/cse/tjg
bsd-05.soe.ucsc.edu:/export/admin	5242880	2761557	2481324	53%	/soe/admin
bsd-05.soe.ucsc.edu:/export/software/old/linuxSW	1048576	744812	2903765	72%	/projects/linuxSW
bsd-05.soe.ucsc.edu:/export/home/zheng6	5242880	2276533	2966347	44%	/cse/zheng6
bsd-05.soe.ucsc.edu:/export/home/compton	5242880	2276533	2966347	44%	/cse/compton
bsd-05.soe.ucsc.edu:/export/home/coco	5242880	2276533	2966347	44%	/cse/coco
bsd-05.soe.ucsc.edu:/export/home/qchen63	5242880	2276533	2966347	44%	/cse/qchen63
bsd-05.soe.ucsc.edu:/export/home/memory	5242880	2276533	2966347	44%	/cse/memory
bsd-05.soe.ucsc.edu:/export/projects/lings	6891456	4971128	1378329	79%	/projects/lings
bsd-05.soe.ucsc.edu:/export/home/musher	5242880	2276533	2966347	44%	/cse/musher

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.


17

## Kernel File System Structures

- In-memory **mount table**
- **Directory cache**
- **Buffers** for reading and writing
  - Usually circular buffers
- **System wide open file table:**
  - File control blocks
  - Data buffers
- **Per-process open file tables:**
  - All files that are open as a result of an `open()` system call not matched with a corresponding `close()`
  - You are doing this in Lab 3 system calls section!
  - Position within the file as a result of the `seek()` or `read()` system calls

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

18



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved. <http://lectures.csail.mit.edu/2013.06.1/lab3/file-system-hierarchy.pdf>

19

## Mounting a New File System


- Consider a Unix computer with one disk, formatted as a single partition mounted as the root directory "/"
- Users' home directories are beneath /home  
`/home/fred`, `/home/barney`, `/home/wilma`, `/home/betty` etc.
- Some time later, the single file system is getting full with user junk so the system administrator decides to add a new HDD just for the users, they:
  - Partition the disk into a single file system and format it
  - Mount the new file system on `/home`
  - Problem?
- Instead
  - Mount new file system on `/mnt`
  - Recursively copy all user directories from `/home` to `/mnt`
  - Recursively remove all directories below `/home`
  - Un-mount new file system from `/mnt`
  - Mount the new file system on `/home`

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

20

## Device Directory

- For each file in a **partition**:
  - Name
  - Type (file, directory, link)
  - Address
  - Current length
  - Maximum length
  - Date last accessed
  - Date last updated
  - Owner ID
  - Protection information



Multiple partitions  
on one disk

One partition on  
multiple disks

21

## Device Directory Operations

- For any given file:
  - Search by name
  - **Create**
  - **Write**
  - **Read**
  - **Delete**
  - **Rename**
- List a directory
- Traverse the file system (enumerate its contents)

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

22


## Device Directory Fundamentals

- Goals:
  - **Efficiency**
    - Locate files quickly
  - **Naming**
    - Multiple users can have same name for multiple files
    - One file can have multiple different names
  - **Grouping**
    - Logical grouping of files by properties
- Possible Directory Structures:
  - Single Level
  - Two Level
  - Acyclic Graph (Tree)

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

23

## Single Level (“Flat”) Directory Structure




**Problems?** Naming & Grouping ☺

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

24

## Two-Level Directory Structure




**Problems?** Grouping ☺

**Benefits?** Path + Name => Simple & Efficient ☺

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

25

## Graph Directory Structure & Links



26

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

## Hard & Symbolic Links

- **Hard links** place a single physical file ( a file on disk ) in multiple locations in the ( logical ) directory graph

```
$ ln <original-file-path> <new-link-path>
$ ln /home/fred/laundrylist.xls fredslaundrylist.xls
$ ls -l *.xls
lrwxr-xr-x 1 david staff 1007 22 Nov 10:11 fredslaundrylist.xls
```

- **Symbolic links** make files appear as if they are in the directory graph in multiple locations

```
$ ln -s <real-file-path> <symbolic-link-path>
$ ln -s /home/fred/pintos mypintos
$ ls -l mypintos
lrwxr-xr-x 1 david staff 8 22 Nov 10:12 mypintos -> /home/fred/pintos
```

- Consider a graph structured directory with a **file X** and two symbolic links to X (**sym1** and **sym2**)
  - What happens if we delete file X but not links sym1 and sym2?
  - **How you might avoid this problem?**

27

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

## Directory Implementation

- **Linear List**
  - Simple to program but time-consuming to execute
    - Easy: Add
    - Hard: Delete, Find
- **Hash Table**
  - A **linear list** stores directory entries
  - A **hash table** takes a value computed from file name and returns a pointer to one location in the list
  - **Problems?**
    - Have to be careful with **name hash collisions**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

28

## File Consistency Semantics

- What is **file consistency**?
  - How multiple simultaneous users access a shared file
- **UNIX File System (UFS)**
  - Modifications immediately visible
  - Can share Location Pointer
    - All sharing processes progress through the file together
- **Andrew File System (AFS)**
  - A **Distributed File System**
  - Modifications not immediately visible
    - Once file is closed, changes become visible to future sessions, but existing sessions do not see the changes unless they close and reopen the file
    - When you login to noggin, nogbad or unix your files are in AFS


29

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.



## Network Attached Storage

- Network-attached storage (NAS) is disk space made available over a network rather than over a local connection (such as a system bus)
- **NFS** (Unix) and **SMB** (Windows) are established protocols
- Amazon **S3** is a modern “cloud” NAS implementation
- Implemented via **remote procedure calls** (RPCs) between clients and storage hosts




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

30

## Lab 3 Secret Sauce

## The Per Process Open File Table



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

32



## Pintos Per-Process Open File Table

- **Pintos System Calls:**

```

int open(const char *file); // returns a file descriptor
int read(int fd, void *buffer, unsigned length);
int write(int fd, const void *buffer, unsigned length);
  
```
- **Pintos File System:**

```

struct file *filesys_open(const char *name);
off_t file_read(struct file *, void *, off_t);
off_t file_write(struct file *, const void *, off_t);
  
```
- **You need to:**
  - Map numeric file descriptors to file pointers
  - Cannot use file descriptors 0, 1, or 2 (`stdin`, `stdout`, and `stderr`)
  - Where are you going to store the per process open file table?
  - No need to store access rights in your OFT

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

33

## Next Lecture

- The UNIX File System
- The Network File System
- Mass Storage

34

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.