

PRINCIPLES OF COMPUTER SYSTEMS DESIGN

CSE130

Winter 2020

File Systems III - Allocation & Free Space



Notices

- Assignment 5 will be available **Sunday March 8**
- Week 10 Lectures (Final Prep) will not be webcast
- Assignments 1 to 4 - Grades available this week
- Late submissions are ignored by the automated grading system
 - If you e-mailed me an extension request and it was granted, see me in Week 10 for a manually assigned grade

2

Today's Lecture

- Open File Tables
- Space Management
- Allocation Algorithms
- The Network File System
- Introduction to Assignment 5

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

3

Everything is a File (descriptor)

- **Everything** in Unix and Unix-Like Operating Systems is exposed via the **file system name space**
- File type examples
 - File
 - Directory
 - Symbolic Link
 - Named Pipe
 - Network Socket
 - Device
- An elegant, simple, unified approach; access rights are easy to understand and implement
- When a file is opened by name (via the `open` system call), a **file descriptor** is returned to the user program
- All subsequent access to the file is done by passing the file descriptor to system calls like `read`, `write`, and `close`

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

4

1

System Call : create (in UNIX "creat")

- Application program calls the logical file system

- Logical file system:

- Allocates a new **File Control Block** (FCB)
- Reads appropriate directory into memory
- Updates it with new file name and FCB
- Writes directory back to disk

- Open File Table (OFT)**

- One for whole system
- One per process (not per thread)


File Control Block : Example Fields	
Permissions: owner, group, ACL	
Dates (creation, last access, last write)	
Size	
Disk location(s)	
etc.	
etc.	

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

5

System Call : open

Program: `fd = open(fname, access_mode);`




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

6

Open File Tables

- Repeated searches of directory expensive
- Cache entry to reduce cost
- Multiple Processes can have the same file open**




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

7

System Call : read

Program: `read(fd, ...);`




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

8

2

System Call : close

Program: close(fd);



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

9

Space Management

- How do we allocate chunks of disk space to files so that we make best use of all available space?
 - How do we provide good **performance** and **reliability**?
 - Allocation methods refer to how disk blocks are allocated for files:
 - Contiguous?
 - Linked?
 - Indexed?
 - **Free Space Management**


Bitmap
11001111
00011110
01101111
00011011

A horizontal bar chart illustrating the distribution of blocks on disk. The x-axis is labeled "blocks on disk" and has tick marks at 0 and N. The bars are colored orange and light green, representing different data types or states.

10

Contiguous Allocation

- File defined by base (start block) and length
 - Sequential and direct access supported
 - Difficult to allocate space or increase file size ☹




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

11

Linked Allocation

- File defined by first and last
 - Solves storage problem - any free block will do
 - Direct access not supported (efficiently) Ⓢ




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

12

FAT : File Allocation Table (Windows)

- Variation of **Linked Allocation**
 - Small array of block numbers linked from start to “eof” (end of file)
- Space set aside at the start of each disk volume for the FAT
- Very efficient if FAT is cached (but synchronisation with disk copy is a problem ☺)
- Follow links in the table to find all the blocks allocated to a file




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

13

Indexed Allocation

- File defined by index
- Supports sequential and direct access
- Solves large file storage problem by introducing the overhead of index blocks ☺⊗



Problems?

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

14


Indexed Allocation

- Need index table
 - What size should the index be?
 - Small is preferred to lower overhead, but
 - Cannot index large files
 - Can dynamically chain index blocks to extend the table
 - Requires following the last entry to next index table
- Could try multilevel index, or
- Combined scheme (index + links)

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

15


Multi-Level Allocation



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

16

The i-node




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

17

Example : The Unix File System

- Boot block
 - Operating System Bootstrap
- Super block
 - Static partition information
- Cylinder block
 - Free block bit map(s)
 - Allocation statistics
- i-nodes
 - Fixed number
 - Run out of i-nodes and the file system is "full", even if the disk is not ☺
- Data blocks
 - Raw data of user files




UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

18

UNIX FS : Free Space Management

- **Free Block Bit Map**
 - Requires additional space. E.g.
block size = 2^{12} bytes (4KB)
disk size = 2^{30} bytes (1GB)
no. of blocks = bitmap size =
 $2^{30} / 2^{12} = 2^{18}$ bits = 32KB
- **Free List** (a linked list)
 - Cannot easily get a sense of contiguous space ☺
 - But less wasted space ☺
- **Counting List**
 - A **disk block address** and a **count**
 - Reduces the length of the **free list** ☺



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

19

UNIX FS : Efficiency & Performance

- Efficiency dependent on:
 - Disk allocation and directory algorithms
 - Types and size of data kept in file's directory entry
- Enhanced Performance
 - **Disk Cache**
 - Separate section of main memory for frequently used blocks
 - **Read-Ahead and Free-Behind**
 - If it appears user code is reading a large file sequentially (a very common action)
 - **Read into cache the next few disk blocks, even if they haven't been asked for yet**
 - **Also remove from cache the blocks that have already been dealt with, even if the user code has not explicitly let go of them by closing the file**
 - **Seek Scheduling**
 - Covered on later slides

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

20

File Systems : Self Study

- Tanenbaum & Bos : Section 10.6 - The Linux File System

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

21

NFS : The Network File System

- Interconnected machines sharing file systems in a transparent manner (via UDP)
 - A **remote** directory is mounted **locally**
 - The mounted directory looks like an integral sub-tree of the local file system, replacing the sub-tree descending from the local directory
 - Specification of the remote directory for the mount operation is **nontransparent**; the host name of the remote directory has to be provided
 - Files in the remote directory can then be accessed in a transparent manner
 - Subject to access-rights accreditation, potentially any file system (or directory within a file system) can be mounted remotely on top of any local directory**

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

22


NFS : The Network File System

- NFS is designed to operate in a **heterogeneous environment** of different machines, operating systems, and network architectures; the NFS specification is platform independent
- This independence is achieved through the use of RPC primitives built on top of an **External Data Representation (XDR)** protocol used between two implementation-independent interfaces
- The NFS specification includes a mount protocol and a separate network protocol for the remote-file-access services

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

23

NFS : The Network File System



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

24

NFS : Protocol Summary

- Provides a set of **remote procedure calls** for remote file operations:
 - Searching for a file within a directory
 - Reading a set of directory entries
 - Manipulating links and directories
 - Accessing file attributes
 - Reading and writing files
- Servers up to Version 3 are **stateless**
 - Each request must provide a full set of arguments
- Version 4 servers are **statefull**
- Modified data must be committed to the server's disk before results are returned to the client (lose advantages of caching)
- Does **not provide concurrency-control** mechanisms

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

25

NFS : Remote Operations

- File-Block Cache**
 - On opening a remote file, kernel checks with the remote server whether to fetch or revalidate cached attributes
 - Cached file blocks used only if its cached attributes are up to date
- File-Attribute Cache**
 - Attribute cache updated when new attributes arrive from the server
 - Clients do not free delayed-write blocks until the server confirms that the data has been written

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

26

File Systems : Summary


- Virtual File System
- Mount Points
- Device Directory Goals
- System Calls & File Descriptors
- Open File Tables
- Space Management
- File Allocation Table (FAT)
- i-nodes
- UNIX File System
- Linux File System (self study)
- Network File System

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

27

Assignment 5

- Basic** (30%)
 - Read & Write ASCII text files
- Advanced** (30%)
 - Append to ASCII text files
 - Copy ASCII text files
- Stretch** (30%)
 - Print a tree of the directory structure



UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.

Do NOT use system() system call
Do NOT use exec() family of system calls
Clearly credit any "borrowed" code

- Especially for Stretch requirement:
- Take your time
 - Start simple
 - Build up solution one-step-at-a-time

28

Next Lecture

- Swap Space
- Mass Storage
- Disk Scheduling

29

UCSC BSOE CSE130 Winter 2020. Copyright © 2017-2020 David C. Harrison. All Rights Reserved.